Cascaded Convolutional Neural Network for Solar Panel Defect Classification

21CSC402P-PROJECT

Submitted by

Divyanshi Jain [RA2311026010639]
Syed Mohammed Ghouse S [RA2311026010643]
Purvee Singh [RA2311026010645]

Under the Guidance of

Dr. Illakiya T
(Assistant Professor, Department of Computational Intelligence)

in partial fulfillment of the requirements for the degree of

BACHELOR OF TECHNOLOGY
in
COMPUTER SCIENCE ENGINEERING
with specialization in Artificial Intelligence and Machine Learning



DEPARTMENT OF COMPUTATIONAL INTELLIGENCE COLLEGE OF ENGINEERING AND TECHNOLOGY
SRM INSTITUTE OF SCIENCE AND TECHNOLOGY
KATTANKULATHUR- 603 203
NOVEMBER 2026

Department of Computational Intelligence

SRM Institute of Science & Technology

Own Work Declaration Form

To be completed by the student for all assessments

Degree/ Course	**: B. Tech Computer Science with specialization in AI/ML**
Student Name	**: Divyanshi Jain, Syed Mohammed Ghouse S, Purvee Singh**
Registration Number	**: RA2311026010639, RA2311026010643, RA2311026010645**
Title of Work	**: Cascaded Convolutional Neural Network for Solar Panel Defect Classification**

We hereby certify that this assessment compiles with the University's Rules and Regulations relating to Academic misconduct and plagiarism, as listed in the University Website, Regulations, and the Education Committee guidelines.

We confirm that all the work contained in this assessment is our own except where indicated, and that I have met the following conditions:

- Clearly referenced / listed all sources as appropriate
- Referenced and put in inverted commas all quoted text (from books, web, etc.)
- Given the sources of all pictures, data etc. that are not my own
- Not made any use of the report(s) or essay(s) of any other student(s) either past or present
- Acknowledged in appropriate places any help that I have received from others (e.g.fellow students, technicians, statisticians, external sources)
- Compiled with any other plagiarism criteria specified in the Course handbook /University website

We understand that any false claim for this work will be penalized in accordance with the University policies and regulations.

DECLARATION:

We are aware of and understand the University's policy on Academic misconduct and plagiarism and we certify that this assessment is our own work, except were indicated by referring, and that we have followed the good academic practices noted above.

Divyanshi Jain, RA2311026010639
Syed Mohammed Ghouse S, RA2311026010643
Purvee Singh, RA2311026010645

SRM INSTITUTE OF SCIENCE AND TECHNOLOGY

KATTANKULATHUR – 603 203

BONAFIDE CERTIFICATE

Certified that 21CSC402P - Project report titled “Cascaded Convolutional Neural Network for Solar Panel Defect Classification” is the bonafide work of “Divyanshi Jain [RA2311026010639], Syed Mohammed Ghouse S [RA2311026010643], Purvee Singh [RA2311026010645]” who carried out the project work under my supervision. Certified further, that to the best of my knowledge the work reported herein does not form any other project report or dissertation on the basis of which a degree or award was conferred on an earlier occasion on this or any other candidate.

SIGNATURE

DR. ILLAKIYA T
SUPERVISOR
ASSISTANT PROFESSOR
DEPARTMENT OF
COMPUTATIONAL INTELLIGENCE

SIGNATURE

DR. R. ANNIE UTHRA
PROFESSOR &HEAD
DEPARTMENT OF
COMPUTATIONAL INTELLIGENCE

ACKNOWLEDGEMENT

We express our humble gratitude to **Dr. C. Muthamizhchelvan**, Vice-Chancellor, SRM Institute of Science and Technology, for the facilities extended for the project work and his continued support.

We extend our sincere thanks to **Dr. Leenus Jesu Martin M,** Dean-CET, SRM Institute of Science and Technology, for his invaluable support.

We wish to thank **Dr. Revathi Venkataraman**, Professor and Chairperson, School of Computing, SRM Institute of Science and Technology, for her support throughout the project work.

We encompass our sincere thanks to, **Dr. M. Pushpalatha**, Professor and Associate Chairperson - CS, School of Computing and **Dr. C. Lakshmi,** Professor and Associate Chairperson -AI, School of Computing, SRM Institute of Science and Technology, for their invaluable support.

We are incredibly grateful to our Head of the Department, **Dr. R. Annie Uthra**, Professor & Department of Computational Intelligence, SRM Institute of Science and Technology, for her suggestions and encouragement at all the stages of the project work.

We want to convey our thanks to our Project Coordinators, Panel Head, and Panel Members Department of Computational Intelligence, SRM Institute of Science and Technology, for their inputs during the project reviews and support.

We register our immeasurable thanks to our Faculty Advisor, Dr. Illakiya T, Department of Computational Intelligence, SRM Institute of Science and Technology, for leading and helping us to complete our course.

Our inexpressible respect and thanks to our guide, Dr. Illakiya T, Department of Computational Intelligence, SRM Institute of Science and Technology, for providing us with an opportunity to pursue our project under their mentorship, and for the freedom and support to explore the topics of our interest.

We sincerely thank all the staff members of Department of Computational Intelligence, School of Computing, S.R.M Institute of Science and Technology, for their help during our project. Finally, we would like to thank our parents, family members, and friends for their unconditional love, constant support and encouragement.

Divyanshi Jain
Syed Mohammed Ghouse S
Purvee Singh

ABSTRACT

Automated inspection of large photovoltaic (PV) installations is essential for sustained power output and safety, but manual thermal surveys do not scale and flat single-stage classifiers struggle with the severe class imbalance and look-alike faults of the InfraredSolarModules benchmark. This project presents Cascaded-SolarNet, a three-stage cascade of five small convolutional neural networks that classifies the 12 classes of the benchmark from single-modality infrared images of only 24 x 40 pixels. Stage 1 separates healthy from defective panels and lets healthy panels exit early; Stage 2 assigns a defect to a fault family; and Stage 3 applies one specialist network per family to name the exact fault. The system was implemented in PyTorch and evaluated on a fixed stratified held-out split of 3,000 images that is never used for training, early stopping or model selection. Removing brightness, contrast and noise augmentation, which destroys the temperature signal of thermal images, raised the from-scratch cascade from 61.7% to 79.8% accuracy, and fine-tuning an ImageNet-pretrained EfficientNet-B0 at every stage reached 83.2% accuracy, macro-F1 0.700 and 95.8% on the healthy-versus-defect decision. A flat classifier with the same backbone is essentially tied (82.6%), so the cascade did not show a clear accuracy advantage, and the family-routing stage remains the largest source of errors. Published 12-class results of up to 93.9% rely on cross-validation and pretrained features and are not directly comparable with a fixed hold-out; the report presents this comparison and its limits. Code, trained models, architecture diagrams and interactive demonstrations are provided.

Index Terms: Solar panel defect classification, infrared thermography, cascaded convolutional neural network, early exit, hierarchical classification, transfer learning, class imbalance.

TABLE OF CONTENTS

CHAPTER NO.	TITLE	PAGE NO.
	ABSTRACT	**vi**
	TABLE OF CONTENTS	**vii**
	LIST OF FIGURES	**xii**
	LIST OF TABLES	**xiv**
	ABBREVIATIONS	**xvi**
1	**INTRODUCTION**	**1**
1.1	Introduction to Solar Panel Defect Classification	1
1.2	Problem Statement and Description	1
1.3	Motivation	2
1.4	Sustainable Development Goal of the Project	3
2	**LITERATURE SURVEY**	**4**
2.1	Overview of PV Defect Detection and Classification Methods	4
2.1.1	Imaging Modalities for PV Inspection	4
2.1.2	The InfraredSolarModules Benchmark	4
2.1.3	Evolution of the Methods	4
2.2	Deep Learning Approaches for PV Defect Analysis	5
2.2.1	Infrared Solar Module Dataset (Millendorf et al., 2020)	5
2.2.2	Binary Cascaded CNN with Thermal-Visual Augmentation (Pathak and Patil, 2025)	5
2.2.3	PV-faultNet: Lightweight Optimised CNN (2024)	6
2.2.4	Defect Detection Based on Improved VarifocalNet (2024)	6
2.2.5	Segmentation of Cell-Level Anomalies in EL Images (Otamendi et al., 2021)	6
2.2.6	Field-Adapted ResNet-50 with Transfer Learning	7

2.2.7	Deep Learning Ensembles for PV Defect Classification (Tella et al., 2025)	7
2.2.8	CNN-Based Classification on RGB Images	7
2.2.9	CNN with Bacterial Foraging Optimisation	8
2.2.10	Cascade R-CNN for Defect Detection	8
2.2.11	Later Work on the InfraredSolarModules Benchmark	8
2.3	Comparative Analysis	9
2.3.1	Accuracy Versus Available Information	10
2.3.2	Binary Versus Multi-Class Classification	10
2.3.3	Evaluation Protocols	10
2.4	Limitations of Existing Methods	10
2.5	Research Gaps and Contributions	11
2.6	Research Objectives	11
2.7	Plan of Action	11
3	**PROBLEM FORMULATION**	**12**
3.1	Formal Problem Definition	12
3.2	Decomposition into Sub-Problems	12
3.2.1	Healthy Versus Defective (Triage)	12
3.2.2	Fault Family	12
3.2.3	Fine-Grained Subtype	13
3.3	Constraints and Assumptions	13
3.3.1	Dataset	13
3.3.2	Input Modality and Image Size	13
3.3.3	Whole-Image Classification	13
3.3.4	Evaluation Protocol	14
3.4	Why the Problem Is Non-Trivial	14
3.4.1	Class Imbalance	14

3.4.2	Look-Alike Faults	14
3.4.3	Error Propagation	14
3.5	Problem Statement Summary	14
4	**PROPOSED SOLUTION**	**15**
4.1	Hierarchical Three-Stage Cascade	15
4.2	Early Exit for Healthy Panels	15
4.3	Fault-Family Grouping	15
4.4	Two Backbones: Lightweight From Scratch and ImageNet-Pretrained	15
4.5	Single-Modality, Whole-Image Pipeline	16
4.6	Summary of Contributions	16
5	**ARCHITECTURE AND IMPLEMENTATION**	**17**
5.1	Overview of the Implemented System	17
5.2	Data Pipeline	18
5.2.1	Dataset and Class Distribution	18
5.2.2	Split and Normalisation	19
5.2.3	Augmentation	19
5.2.4	Handling Class Imbalance	19
5.3	Building Blocks and Their Mathematics	20
5.3.1	Convolution	20
5.3.2	Batch Normalisation	20
5.3.3	ReLU and Max Pooling	20
5.3.4	Global Average Pooling and Classification Head	21
5.3.5	Loss Function	21
5.3.6	Counting Parameters	21
5.4	Stage 1: CNN-1, Binary Triage	21
5.5	Stage 2: CNN-2, Fault-Family Classifier	22

5.6	Stage 3: Fine-Grained Specialists	23
5.7	Pretrained Variant	25
5.8	Training Procedure	25
5.9	Cascade Inference and Cost Analysis	26
5.10	Evaluation Tools and Demonstrations	26
5.11	Testing and Reproducibility	27
5.12	Chapter Summary	28
6	**RESULTS AND ANALYSIS**	**29**
6.1	Experimental Setup	29
6.1.1	Dataset and Protocol	29
6.1.2	Implementation and Training Configuration	29
6.1.3	Evaluation Metrics	29
6.1.4	Configurations	30
6.2	Overall Performance	30
6.2.1	Test Results	30
6.2.2	Per-Network Validation Results and Training Behaviour	31
6.3	Ablation Study	32
6.3.1	Photometric Augmentation	32
6.3.2	Input Size and Network Capacity	33
6.3.3	Balanced Sampling and Label Smoothing	33
6.3.4	Inference-Time Variants	33
6.3.5	Fault-Family Grouping	33
6.3.6	ImageNet Pretraining	34
6.3.7	Summary of the Ablations	34
6.4	Error Analysis by Stage	34
6.5	Per-Class Performance and Confusion	35

6.5.1	Per-Class F1	35
6.5.2	Confusion Matrices	36
6.6	Comparison with Published Work	38
6.7	Data Overlap Analysis	40
6.8	Discussion and Limitations	40
6.8.1	Interpretation of Findings	40
6.8.2	Limitations	40
6.8.3	Threats to Validity	41
6.9	Chapter Summary	41
7	**CONCLUSION AND FUTURE SCOPE**	**42**
7.1	Conclusion	42
7.2	Future Scope	42
	REFERENCES	**43**
A	**CONFIGURATION FILES**	**45**
B	**KEY SOURCE CODE**	**47**
C	**RUNNING THE DEMONSTRATIONS**	**50**
D	**REPOSITORY STRUCTURE**	**52**

LIST OF FIGURES

FIGURE NO.	TITLE	PAGE NO.
1	1. The implemented three-stage Cascaded-SolarNet architecture	17
2	2. Class distribution of the InfraredSolarModules dataset (logarithmic axis)	18
3	3. One test image per class, shown with the inferno colour map	19
4	4. Stage 1: CNN-1, binary triage (layers and feature-map sizes)	22
5	5. Stage 2: CNN-2, fault-family classifier	23
6	6. Stage 3: the three specialists and their classes	24
7	7. Output of the live demonstration: each stage's decision and confidence for Hot-Spot test images	27
8	8. Cascade test accuracy and macro-F1 by configuration (dashed line: from-scratch flat CNN)	31
9	9. Best validation macro-F1 of each stage network, from scratch and pretrained	31
10	10. Validation macro-F1 and training loss of the five stage networks	32
11	11. Where the cascade makes its errors, by configuration	35
12	12. Per-class F1 on the test set	36

FIGURE NO.	TITLE	PAGE NO.
13	13. Confusion matrix of the pretrained cascade on the test set (row-normalised, counts inside cells)	37
14	14. Confusion matrix of the pretrained flat CNN on the test set	38
15	15. Example output of the command-line demonstration for random test images (green: correct, red: wrong)	51

LIST OF TABLES

TABLE NO.	TITLE	PAGE NO.
2.1	Summary of the reviewed studies	9
2.2	Comparison of approach families	9
2.3	Results reported on InfraredSolarModules	9
2.4	Research gaps and the contributions that address them	11
3.1	Class distribution and stratified split (seed 42)	13
4.1	Contributions and the gaps they address	16
5.1	Layer-by-layer structure of CNN-1 (from-scratch variant)	22
5.2	Layer-by-layer structure of CNN-2 (from-scratch variant)	23
5.3	Parameter counts of the five stage networks	24
5.4	Architecture summary of the five networks (from-scratch variant)	24
5.5	Training data of each stage network	25
5.6	Automated tests	27
6.1	Training configuration of the evaluated models	29
6.2	Definition of the evaluated configurations	30
6.3	Test results by configuration	30
6.4	Effect of photometric augmentation (cascade, 40 x 24 input)	32
6.5	Effect of input size and capacity (cascade, flips only)	33
6.6	Inference-time variants on the improved cascade (no retraining)	33
6.7	Semantic versus data-driven family grouping (from-scratch	34

TABLE NO.	TITLE	PAGE NO.
	cascade)	
6.8	Test errors per stage (out of 3,000 images)	34
6.9	Per-class F1 on the test set	35
6.10	Most frequent confusions of the pretrained cascade	38
6.11	Published results on InfraredSolarModules and results of this project	39
6.12	Similarity between test and training images (raw pixels)	40
D.1	Layout of the project repository	52

ABBREVIATIONS

ABBREVIATION	FULL FORM
AI	Artificial Intelligence
AMP	Automatic Mixed Precision
BN	Batch Normalisation
CNN	Convolutional Neural Network
CV	Cross-Validation
EL	Electroluminescence
F1	Harmonic mean of precision and recall
FC	Fully Connected (layer)
GAP	Global Average Pooling
GPU	Graphics Processing Unit
IR	Infrared
NCA	Neighbourhood Component Analysis
PV	Photovoltaic
ReLU	Rectified Linear Unit
RF	Random Forest
SDG	Sustainable Development Goal
SVM	Support Vector Machine
TTA	Test-Time Augmentation
UAV	Unmanned Aerial Vehicle
ViT	Vision Transformer

CHAPTER 1

INTRODUCTION

Solar farms contain thousands of photovoltaic modules that develop defects over time, including hot spots, cracks, soiling, shadowing and diode failures. Each of these reduces power output, and some, such as undetected hot spots, create a fire risk. Manual thermal inspection is slow, expensive, labour-intensive and error-prone, and it cannot scale to utility-grade installations. Deep learning applied to drone-captured infrared (IR) imagery offers an automated alternative: an unmanned aerial vehicle with a thermal camera can image thousands of modules in one flight, and a trained classifier can label each module in milliseconds.

This project builds and evaluates such a classifier on the InfraredSolarModules benchmark, a public dataset of 20,000 IR images of 24 x 40 pixels with 12 classes (one healthy class and eleven fault types). Instead of a single network that chooses directly among all twelve classes, the project proposes a cascade of small networks that answers three easier questions in turn.

1.1 Introduction to Solar Panel Defect Classification

In an infrared image brightness encodes temperature. A healthy module heats fairly uniformly during operation, whereas a defect produces a localised temperature anomaly: a hot spot, a cold region or a characteristic thermal pattern. Classifying these patterns is therefore a pattern-recognition problem on very small images, where fine texture is absent and the position, size and shape of heat concentrations carry the information. Typical defect families are electrical faults (cells and bypass diodes), environmental conditions (soiling, shadows, vegetation) and structural damage (cracks, offline modules).

The task is difficult for three reasons. The classes are strongly imbalanced, because half of the images show healthy modules. Several fault types look alike at low resolution. And the system has to be cheap enough to run on drone hardware. These properties motivate a design that removes the dominant class first and then specialises.

1.2 Problem Statement and Description

Most existing CNN approaches use a flat, single-stage classifier that must choose directly among all classes. This design has two weaknesses. First, visually similar fault types produce overlapping features, so the network confuses them. Second, healthy panels vastly outnumber faulty ones, which biases the network towards the majority class and hurts the rare faults. The

literature reviewed in this project reports that accuracy falls from about 73% for binary tasks to about 68% for multi-class tasks across eight architectures, and that single-stage results on this benchmark stop at about 93%.

The problem addressed here is to classify a single-modality 24 x 40 pixel IR image of a PV module into one of 12 classes accurately, with a model small enough for edge use, and to understand where the errors of the proposed design come from. The literature survey identified six gaps:

- **G1.** Dual-sensor approaches are impractical. Models reaching 98% to 99% accuracy need paired thermal and visual data that a single-camera drone cannot supply.
- **G2.** Scaling the model gives negligible gains. Heavier backbones add little accuracy over lightweight ones.
- **G3.** Bounding-box detectors split thin defects. Detection models mislocalise fine crack boundaries.
- **G4.** Naive cascades add latency. Detection-plus-classification pipelines are too slow for real-time use.
- **G5.** Accuracy drops as the number of classes grows, from roughly 73% for binary to 68% for multi-class classification.
- **G6.** Similar defect types are systematically confused, and the confusion persists in flat classifiers.

1.3 Motivation

The practical cost of misclassification is concrete. A missed hot spot can become a fire; a false alarm sends a maintenance crew to a healthy module. Because most modules are healthy, a classifier that wastes effort on them is both slow and biased. A hierarchical design addresses this in a natural way: a small first network handles the easy and most frequent decision and passes only suspicious modules on, and later networks can then spend their capacity on the genuinely hard distinctions.

A second motivation is scientific. Cascades are often assumed to beat flat classifiers, but the literature offers few controlled comparisons on a small thermal benchmark. This project measures that claim with identical building blocks, identical data splits and a stage-by-stage attribution of errors, and reports the result honestly, including where the cascade does not help.

1.4 Sustainable Development Goal of the Project

The project contributes most directly to SDG 7, Affordable and Clean Energy. Solar installations that are inspected faster and more cheaply lose less energy to undetected defects, which improves the return on clean-energy investment. It also supports SDG 13, Climate Action, because every percentage point of recovered solar output displaces fossil generation, and SDG 9, Industry, Innovation and Infrastructure, through low-cost automated inspection that can run on small drones.

CHAPTER 2

LITERATURE SURVEY

This chapter reviews deep-learning approaches to PV defect detection and classification, with emphasis on infrared imagery and on the InfraredSolarModules benchmark, and derives the research gaps and objectives of the project. Ten studies from 2020 to 2025 are discussed in turn, followed by a comparative analysis. The studies were re-checked against their primary sources while preparing the final report; where a result had been attributed to the wrong paper in the first version of the survey, it has been corrected below.

2.1 Overview of PV Defect Detection and Classification Methods

2.1.1 Imaging Modalities for PV Inspection

Three imaging modalities are used for PV inspection. Infrared (IR) thermography is the primary modality for operating plants: in a healthy panel heat is distributed uniformly, whereas defects create localised temperature anomalies such as hot spots, cold regions or characteristic thermal signatures, and a drone can capture thousands of thermal images per flight. Electroluminescence (EL) imaging reveals cell-level cracks and inactive regions with very high detail, but requires the panel to be electrically driven and imaged in darkness, so it suits factory and laboratory inspection and not routine drone surveys. Visible-spectrum (RGB) imaging is cheap and detailed but shows only what is visible on the surface, such as dust, shadows and gross cracks, and says little about electrical faults.

2.1.2 The InfraredSolarModules Benchmark

The InfraredSolarModules dataset, released by Raptor Maps [1], standardises 12 classes from real commercial drone flights: 11 anomaly types and one No-Anomaly class. The images are only 24 x 40 pixels because of the altitude of capture, so classification depends mainly on recognising temperature-pattern signatures and not on structural detail. The classes fall naturally into electrical faults (Cell, Cell-Multi, Diode, Diode-Multi, Hot-Spot, Hot-Spot-Multi), environmental conditions (Soiling, Shadowing, Vegetation) and structural conditions (Cracking, Offline-Module), which motivates hierarchical strategies over flat ones.

2.1.3 Evolution of the Methods

Automated analysis of PV imagery has progressed through several generations. Early work

extracted handcrafted thermal features and classified them with classical learners such as random forests. The second generation applied convolutional networks end to end, from small custom CNNs to ImageNet-pretrained VGG, ResNet and EfficientNet backbones, usually with transfer learning because labelled PV datasets are small. A third thread moved to object detection and segmentation (Faster R-CNN, Cascade R-CNN, VarifocalNet, YOLO) to localise defects within a module or panel image. The most recent work explores cascades, ensembles, feature-selection pipelines and multi-sensor fusion, usually at the cost of extra sensors or computation.

2.2 Deep Learning Approaches for PV Defect Analysis

2.2.1 Infrared Solar Module Dataset (Millendorf et al., 2020)

Millendorf, Obropta and Vadhavkar of Raptor Maps introduced the benchmark at the ICLR 2020 Workshop on AI for Earth Sciences [1]. The paper motivates automation with industry figures: aerial inspections find anomalies in a few percent of modules, and the dataset was released to let researchers quickly train models and tune hyper-parameters. It describes the 12 classes and the data format and lists benchmarking of classification results as a goal, but it reports no model accuracies. The widely quoted baselines on this dataset (Random Forest 84%, custom CNN 92.5%, EfficientNet with an SVM 93%) were published in later papers, which are discussed in Section 2.2.11.

2.2.2 Binary Cascaded CNN with Thermal-Visual Augmentation (Pathak and Patil, 2025)

Pathak and Patil published a cascaded CNN in the International Journal of Image, Graphics and Signal Processing (vol. 17, no. 4, 2025) that fuses thermal infrared images with co-registered visible-spectrum images [5]. Their Binary Cascaded CNN (B-CNN) has two stages: the first classifies a module as healthy or faulty and the second identifies the specific fault. The thermal-visual augmentation greatly enriches the feature space and compensates for the low resolution of IR-only inputs, reaching 99.8% precision and 98.5% accuracy, the highest results in this survey. The critical limitation is the requirement for paired dual-sensor data, which most UAVs, carrying only a thermal camera, cannot provide, and the fusion adds hardware cost, synchronisation effort and computation. This is the closest relative of our cascade, and it shows that a binary-first cascade is effective; our design keeps the cascade idea but uses a single sensor.

2.2.3 PV-faultNet: Lightweight Optimised CNN (2024)

PV-faultNet investigates the design space of CNN architectures for efficient PV defect classification [6]. The authors benchmark architectures from compact custom CNNs to larger Siamese networks and multi-spectral models, and find that a carefully designed compact CNN reaches 93% accuracy at 8 ms per image, competitive with much larger models. The key finding is that larger backbones give negligible gains because of the low information content of small IR images: the bottleneck is the input resolution, not the model capacity. The work stresses efficient architecture design for real-time edge inference, but it does not address the accuracy ceiling or class imbalance and remains at about 93% single-stage accuracy. It supports our choice of small per-stage networks.

2.2.4 Defect Detection Based on Improved VarifocalNet (2024)

This study, published in Scientific Reports (vol. 14, 2024), extends the VarifocalNet detector to localise defects in drone-captured IR images and compares it with Cascade R-CNN and YOLOv7 on a custom dataset of annotated images [7]. VarifocalNet achieves the highest confidence for coarse defect detection, and Cascade R-CNN performs well on common categories, but both struggle with fine crack localisation and produce fragmented or offset bounding boxes for thin linear defects. Detection-based approaches thus add a failure mode beyond classification errors: a defect may be correctly classified but poorly localised, which defeats the inspection purpose. This motivates whole-image classification for the low-resolution IR domain.

2.2.5 Segmentation of Cell-Level Anomalies in EL Images (Otamendi et al., 2021)

Otamendi, Martinez and Quartulli published a pipeline in Solar Energy (vol. 220, pp. 914-926, 2021) for cell-level anomaly detection in electroluminescence images [8]. Faster R-CNN proposes regions and an EfficientNet-based classifier identifies the anomaly type inside each region, reaching high precision on their EL dataset. EL imaging, however, needs electrically active panels imaged in darkness and is much less practical than IR thermography for routine drone inspection, and the detection-then-classification pipeline adds seconds of latency per image, which is incompatible with real-time use. The work nevertheless shows that cell-level granularity is feasible.

2.2.6 Field-Adapted ResNet-50 with Transfer Learning

This study addresses the gap between laboratory PV defect datasets and field conditions. A ResNet-50 pretrained on ImageNet is fine-tuned on field-collected images with extensive augmentation (8x, including rotations, flips, brightness perturbation and noise injection), reaching 98.59% accuracy on the field data. The augmentation is critical because field IR images vary in thermal contrast, panel orientation, atmospheric conditions and camera angle. Limitations are the high resource demand of ResNet-50, which makes edge deployment hard, and that class imbalance and fine-grained subtypes are not addressed. Notably, its brightness and noise augmentation is the kind that our own ablation found harmful on the InfraredSolarModules benchmark, which suggests that the right augmentation depends on how the dataset was captured and normalised.

2.2.7 Deep Learning Ensembles for PV Defect Classification (Tella et al., 2025)

Tella, Hussein and Rehman published a multi-model benchmark in Case Studies in Thermal Engineering (vol. 66, 2025) [9]. It evaluates eight architectures, including VGG-16, ResNet-50, DenseNet, MobileNet, EfficientNet and Vision Transformer variants, on the ELPV benchmark. The key finding is a systematic gap between binary classification (about 73% on average) and multi-class fault classification (about 68% on average) that no single model closes, and the gap reflects the structural difficulty of multi-class decision boundaries between look-alike faults at low resolution and not model capacity. An ensemble gives only marginal gains. This is the main published evidence for the weakness of the flat single-stage paradigm that motivates our decomposition.

2.2.8 CNN-Based Classification on RGB Images

This work applies VGG and ResNet variants to RGB images of solar panels with a four-class taxonomy (normal, cracked, dusty, shadowed). Accuracy is high, about 96% to 98%, because there are few classes and the images are rich. The coarse formulation, however, hides the real operational challenge: distinguishing electrical fault subtypes such as hot spots, diode failures and bypass faults requires thermal imaging and finer granularity. The results cannot be compared with IR studies, and the study is useful mainly as a baseline for RGB-based inspection.

2.2.9 CNN with Bacterial Foraging Optimisation

This study combines five pretrained CNNs (VGG-16, AlexNet, SqueezeNet, GoogleNet and ResNet) with Bacterial Foraging Optimisation (BFO), a bio-inspired metaheuristic, to select the most discriminative subset of features from the concatenated feature vectors. VGG-16 with BFO reaches the best accuracy of 98.75% on their data. The price is a large overhead: the feature selection adds computation in training and inference, and all five backbones must run together, which is impractical for real-time drone hardware. It shows that feature selection over multiple backbones can raise accuracy at an unacceptable latency cost.

2.2.10 Cascade R-CNN for Defect Detection

Cascade R-CNN, proposed by Cai and Vasconcelos [10], is a multi-stage detector in which cascaded heads refine region proposals at progressively higher intersection-over-union thresholds. Applied to PV defects it improves localisation for defects of variable aspect ratio, but it shares the crack-fragmentation problem of the other detectors on thin linear defects. Conceptually, cascade detection operates on bounding-box proposals while a cascaded classifier operates on the class decision space. The work motivates the use of progressive, refined decisions, but argues for applying them at the classification level in the low-resolution IR setting.

2.2.11 Later Work on the InfraredSolarModules Benchmark

Several later papers report results on the benchmark, and they must be read with their protocols in mind. Alves et al. trained a custom CNN with data augmentation and report 92.5% test accuracy with cross-validation for the detection of anomalies (a binary task) and 78.85% for eight selected classes [3]. Duranay extracts features with an ImageNet-pretrained EfficientNet-B0, selects 17,000 of them with Neighbourhood Component Analysis and classifies them with an SVM, reporting 93.93% accuracy and 89.82% F1 on all 12 classes with 10-fold cross-validation [2]. Corley et al. report a random forest baseline of 84.0% and cite further results (a cascading ensemble with a voting CNN at 88% [15] and a NASNet-LSTM model at 85% [16]), and they argue that the dataset may be too simple for deployment at scale because the crop level lacks the context of neighbouring panels [4]. These are the numbers against which Chapter 6 places the results of this project.

2.3 Comparative Analysis

Table 2.1 summarises the studies. Table 2.2 groups the approaches by strategy and Table 2.3 collects the results reported on InfraredSolarModules.

Table 2.1: Summary of the reviewed studies

#	Study	Approach	Key result	Limitation
1	Millendorf et al., 2020	Dataset paper (InfraredSolarModules)	No model results reported	Baselines come from later work
2	Pathak and Patil, 2025	Binary cascaded CNN with thermal-visual fusion	99.8% precision, 98.5% accuracy	Needs paired thermal and visual data
3	PV-faultNet, 2024	Compact CNN versus Siamese and multi-spectral	93% at 8 ms per image	Larger backbones gain little
4	Improved VarifocalNet, 2024	Cascade R-CNN, YOLOv7 comparison	High coarse detection confidence	Mislocalises thin cracks
5	Otamendi et al., 2021	Faster R-CNN + EfficientNet on EL images	High precision	Too slow for real-time drones
6	Field-adapted ResNet-50	Transfer learning with 8x augmentation	98.59% on field data	High resource demand
7	Tella et al., 2025	8 CNN/ViT models on ELPV	About 68% multi-class vs 73% binary	Large binary to multi-class drop
8	CNN on RGB images	4-class normal/cracked/dusty/shadowed	96% to 98%	Coarse classes only
9	CNN + Bacterial Foraging Optimisation	Five pretrained CNNs with feature selection	VGG-16 + BFO: 98.75%	Extra tuning, five backbones
10	Duranay, 2023	EfficientNet-B0 features, NCA, SVM	93.93% on 12 classes (10-fold CV)	Cross-validation, pretrained features

Table 2.2: Comparison of approach families

Approach family	Examples	Strengths	Weaknesses
Flat single-stage CNN	Custom CNNs, PV-faultNet, Alves et al.	Simple, fast, small	Plateau near 93%; confuses look-alike faults; biased by imbalance
Pretrained backbone plus classifier	Duranay, ResNet-50, CNN + BFO	High accuracy from transfer learning	Heavy or multi-backbone; protocols often cross-validated
Detection and segmentation	VarifocalNet, Cascade R-CNN, Otamendi et al.	Locate the defect	Fragment thin defects; slow; need box labels
Multi-sensor cascade	Pathak and Patil	Highest accuracy	Needs thermal and visual sensors and fusion
Ensembles	Tella et al., voting CNNs	Marginal accuracy gains	Cost grows with every model
Single-sensor classification cascade (this project)	Cascaded-SolarNet	One IR camera, small networks, early exit, error attribution	Errors propagate between stages

Table 2.3: Results reported on InfraredSolarModules

Source	Method	Task	Reported accuracy
Duranay 2023 [2]	EfficientNet-B0 features, NCA, SVM	12 classes, 10-fold CV	93.93%
Alves et al. 2021 [3]	Custom CNN with augmentation	Binary detection, cross-validation	92.5%
Alves et al. 2021 [3]	Custom CNN with augmentation	8 selected classes, cross-validation	78.85%
Barraz et al. 2025, as cited in	Cascading ensemble, voting	Not stated	88%

Source	Method	Task	Reported accuracy
[4]	CNN		
Gopalakrishnan et al. 2024, as cited in [4]	NASNet + LSTM	Not stated	85%
Corley et al. 2025 [4]	Random Forest	Anomaly classification	84.0%

2.3.1 Accuracy Versus Available Information

The highest accuracies in the survey (98% to 99%) come from methods that use more information than a single IR crop gives: a second sensor, a different modality with higher detail, field-specific training data or an ensemble of backbones. Methods restricted to single IR crops cluster around 93%. The comparison shows that accuracy is bought with sensors and computation, and that the single-sensor setting is a harder and more practical benchmark than many headline figures suggest.

2.3.2 Binary Versus Multi-Class Classification

Binary healthy-versus-faulty detection reaches 92.5% to 97% in several studies, while multi-class results are lower: about 68% across eight architectures in the ELPV benchmark [9] and 78.85% for eight classes in [3]. This supports decomposing the task so that the easy binary decision is made first and by a dedicated network, which is the design of Cascaded-SolarNet.

2.3.3 Evaluation Protocols

Most reported results use k-fold cross-validation, some preprocessing or feature selection before the split is applied, and several datasets or class definitions. Results from different papers are therefore not directly comparable. In this project all results use one fixed held-out split that is never touched during training or selection, and Chapter 6 discusses what this means for the comparison.

2.4 Limitations of Existing Methods

Four limitations recur across the reviewed work. First, the highest accuracies depend on information that a drone does not have (paired visible images, powered-panel electroluminescence) or on heavy computation (ensembles, large backbones, feature selection over several networks). Second, detection pipelines add failure modes on thin defects and need bounding-box labels that the crop-level benchmark lacks. Third, results are reported under different protocols, mostly cross-validation, which makes them hard to compare, and near-duplicate crops in the dataset can inflate cross-validated scores. Fourth, almost no study reports

where a hierarchical model loses accuracy, so the benefit of cascading remains unmeasured.

2.5 Research Gaps and Contributions

The survey identifies six research gaps. Table 2.4 maps each gap to the contribution of this project.

Table 2.4: Research gaps and the contributions that address them

Gap	Research gap	Contribution
G1	Dual-sensor approaches are impractical	C4: single-modality design using only single-camera IR images
G2	Scaling the model gives negligible gains	C1: five small per-stage networks instead of one large model
G3	Bounding boxes fragment thin defects	C5: whole-image classification with no localisation step
G4	Naive cascades add latency	C2: early-exit binary front end for healthy panels
G5	Accuracy drops as classes grow	C1: hierarchical decomposition into 2-, 3- and 2-to-6-way decisions
G6	Similar defects are confused	C3: family grouping before fine classification, semantic and data-driven

2.6 Research Objectives

- **O1.** To design a three-stage cascade of small CNNs for 12-class IR solar module classification from a single sensor.
- **O2.** To implement the system reproducibly in PyTorch with a fixed stratified split, tests and published code.
- **O3.** To measure, with controlled ablations, which design choices (augmentation, input size, capacity, grouping, pretraining) matter.
- **O4.** To attribute the cascade's errors to the stage that caused them and to compare it with a flat classifier and with published results.

2.7 Plan of Action

The work followed six steps: (1) literature survey and gap analysis; (2) dataset preparation with a stratified 70/15/15 split; (3) implementation of the five-network cascade and a flat baseline; (4) controlled ablations of augmentation, input size, network capacity, family grouping and inference variants; (5) fine-tuning of an ImageNet-pretrained backbone at every stage; and (6) evaluation, error attribution, comparison with published work, architecture diagrams, demonstrations and this report.

CHAPTER 3

PROBLEM FORMULATION

This chapter defines the classification problem precisely, decomposes it into sub-problems, states the constraints under which the system is built and explains why a naive solution is not enough.

3.1 Formal Problem Definition

Let x be a single-channel infrared image of one PV module with 40 rows and 24 columns, x in R^(40 x 24) with values scaled to [0, 1]. Let the label set be Y = {y0, y1, ..., y11}, where y0 is No-Anomaly and y1 to y11 are the fault classes. The goal is to learn a function f that maps x to a label so that the expected accuracy and the macro-averaged F1 score over Y are maximised on unseen images, subject to the constraints of Section 3.3.

The label set carries a hierarchy: Y = {y0} U F_E U F_V U F_S, where the families F_E (Electrical), F_V (Environmental) and F_S (Structural) partition the eleven faults. This structure is exploited by the cascade: f(x) = g_fam(x)(x) if the triage function h(x) predicts a defect, and y0 otherwise.

3.2 Decomposition into Sub-Problems

3.2.1 Healthy Versus Defective (Triage)

The first sub-problem is the binary function h(x) in {healthy, defect}. It is the easiest decision in the system and the one that removes the dominant class. An error here is irreversible: a missed defect is never seen by later stages, and a false alarm only costs computation and a wrong label.

3.2.2 Fault Family

For images judged defective, the second sub-problem is the three-way function that selects a family. Faults that share a family should look similar to each other and different from other families, so this decision should be much easier than an eleven-way choice. Section 6 tests how well this holds.

3.2.3 Fine-Grained Subtype

Within a family, a specialist function chooses the exact fault: six classes for Electrical, three for Environmental and two for Structural. Each specialist is trained only on its own family and can spend all its capacity on the distinctions inside it.

3.3 Constraints and Assumptions

3.3.1 Dataset

The InfraredSolarModules dataset contains 20,000 images from real drone flights, 24 pixels wide and 40 pixels tall. The classes and their distribution in the stratified split used throughout this report are given in Table 3.1.

Table 3.1: Class distribution and stratified split (seed 42)

Class	Family	Total	Train	Validation	Test
No-Anomaly	Healthy	10000	7000	1500	1500
Cell	Electrical	1877	1314	282	281
Cell-Multi	Electrical	1288	902	193	193
Diode	Electrical	1499	1049	225	225
Diode-Multi	Electrical	175	123	26	26
Hot-Spot	Electrical	249	174	38	37
Hot-Spot-Multi	Electrical	246	172	37	37
Soiling	Environmental	204	143	30	31
Shadowing	Environmental	1056	739	158	159
Vegetation	Environmental	1639	1147	246	246
Cracking	Structural	940	658	141	141
Offline-Module	Structural	827	579	124	124
Total		**20000**	**14000**	**3000**	**3000**

3.3.2 Input Modality and Image Size

Only single-channel IR images are used, because a single-camera drone cannot supply paired visible images. The images are tiny, so networks must recognise where heat is concentrated and what shape it takes rather than rely on texture.

3.3.3 Whole-Image Classification

The system classifies the whole module image and does not localise defects. Detection pipelines were shown in the literature to fragment thin defects, and the dataset provides one label per crop and no boxes.

3.3.4 Evaluation Protocol

All results are measured on one fixed, stratified held-out test split of 3,000 images. The test images are never used for training, early stopping, normalisation statistics or model selection; all tuning uses a separate 3,000-image validation split.

3.4 Why the Problem Is Non-Trivial

3.4.1 Class Imbalance

Half of the images are healthy, and the rarest classes have 175 to 249 images in total (123 to 174 for training). A classifier that always answers No-Anomaly scores 50% accuracy and learns nothing about faults.

3.4.2 Look-Alike Faults

Several classes differ only in how many cells or diodes are affected (Cell and Cell-Multi, Diode and Diode-Multi, Hot-Spot and Hot-Spot-Multi), and at 24 x 40 pixels these differences are subtle.

3.4.3 Error Propagation

In a cascade, a wrong decision at an early stage cannot be corrected later. The overall accuracy of a defective image is roughly the product of the accuracies of the stages it passes through, so the early stages must be very accurate. Whether a cascade nevertheless beats a flat classifier is an empirical question, which this project answers by measurement.

3.5 Problem Statement Summary

Given single-modality 40 x 24 IR images of PV modules, build a small-network system that classifies them into 12 classes, evaluate it on a clean held-out split, attribute its errors to individual stages and compare it with a flat classifier and with published results.

CHAPTER 4

PROPOSED SOLUTION

This chapter describes the design decisions of Cascaded-SolarNet and relates each one to a gap identified in Chapter 2. The detailed architecture and implementation follow in Chapter 5.

4.1 Hierarchical Three-Stage Cascade

Cascaded-SolarNet replaces one large 12-way classifier with five small specialist CNNs arranged in three stages. Each network answers one small, well-defined question, and the answer of one stage decides which network runs next. A healthy panel passes through only one network; a defective panel passes through exactly three: the triage network, the family network and one specialist.

4.2 Early Exit for Healthy Panels

Stage 1 is a binary classifier. When it predicts No-Anomaly the pipeline outputs that label and stops. Because healthy panels are the large majority, this removes the dominant class from all later networks, handles class imbalance by structure instead of relying only on loss re-weighting, and in principle saves computation on most images.

4.3 Fault-Family Grouping

Stage 2 routes each defect to one of three families. The semantic grouping follows the physical origin of the faults: Electrical (Cell, Cell-Multi, Diode, Diode-Multi, Hot-Spot, Hot-Spot-Multi), Environmental (Soiling, Shadowing, Vegetation) and Structural (Cracking, Offline-Module). As an alternative, a data-driven grouping was derived by testing every partition of the eleven faults into three families of at least three classes and choosing the one that minimises the confusions of a trained flat network on the validation split: group A (Cell, Cell-Multi, Soiling, Vegetation, Cracking), group B (Diode, Diode-Multi, Offline-Module) and group C (Hot-Spot, Hot-Spot-Multi, Shadowing). This reduced cross-family confusions on the validation split from 264 to 87.

4.4 Two Backbones: Lightweight From Scratch and ImageNet-Pretrained

Every stage network is available in two variants that share the cascade logic. The from-scratch

variant uses small convolutional blocks and has no external data. The pretrained variant uses an ImageNet-pretrained EfficientNet-B0 with a new classification head and is fine-tuned on the stage's data. The comparison of the two variants quantifies what pretraining adds on this tiny thermal domain.

4.5 Single-Modality, Whole-Image Pipeline

The system takes only a single-channel IR image and outputs a label for the whole module. It needs no second sensor, no alignment step and no localisation, which removes the failure modes G1 and G3.

4.6 Summary of Contributions

Table 4.1: Contributions and the gaps they address

Contribution	Description	Gaps addressed
C1	Hierarchical three-stage cascade in place of a flat 12-class classifier	G2, G5
C2	Early-exit binary front end that removes the dominant No-Anomaly class	G4, imbalance
C3	Fault-family grouping (semantic and data-driven) before fine classification	G6
C4	Single-modality design using only single-camera IR images	G1
C5	Pure classification pipeline without localisation or bounding boxes	G3
C6	Controlled ablations and per-stage error attribution, with a leak-free evaluation protocol	Comparability

CHAPTER 5

ARCHITECTURE AND IMPLEMENTATION

5.1 Overview of the Implemented System

Figure 1 shows the implemented cascade. The system is implemented in Python 3.12 with PyTorch 2.11 and torchvision, scikit-learn for the metrics and Matplotlib for plots. Training and evaluation ran on an NVIDIA GeForce RTX 5060 Laptop GPU with 8 GB of memory. The code is organised as a Python package (solarnet) with separate modules for the class hierarchy (taxonomy), data handling (data), the models, shared helpers (common), training (train) and evaluation (evaluate), plus configuration files, scripts for the baselines, cross-validation, figure and report generation, and a test suite. Appendix D lists the repository layout.

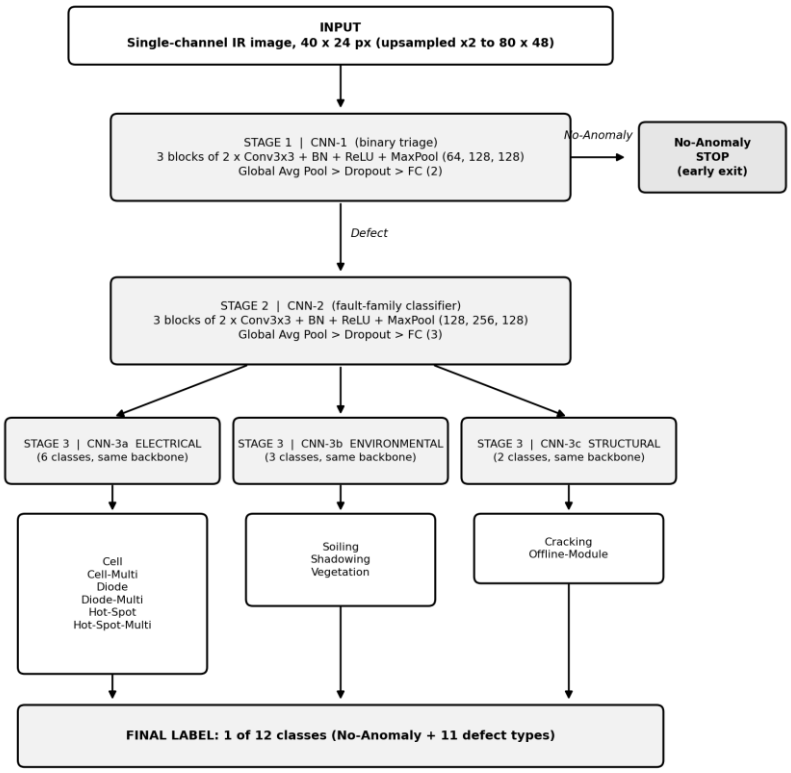


Fig. 1. The implemented three-stage Cascaded-SolarNet architecture.

An image is processed as follows. It is normalised and passed to CNN-1. If CNN-1 predicts No-Anomaly the pipeline stops. Otherwise CNN-2 predicts a family and the image is passed

to the specialist of that family, whose prediction is converted to a global class index by adding the family's class offset. Every network is trained independently on ground-truth routing, and the cascade is assembled only at evaluation time, so the error of a stage is measured exactly as it would occur in deployment.

5.2 Data Pipeline

5.2.1 Dataset and Class Distribution

The InfraredSolarModules dataset has 20,000 grayscale images of 24 pixels wide and 40 pixels tall. Figure 2 shows the class distribution on a logarithmic axis: half of the images are healthy, and the four rarest classes (Diode-Multi, Soiling, Hot-Spot-Multi and Hot-Spot) have between 175 and 249 images each. Figure 3 shows one test image per class with a thermal colour map. Even at this size the healthy panel is nearly uniform, while faults show localised bright regions whose position and extent differ by class.

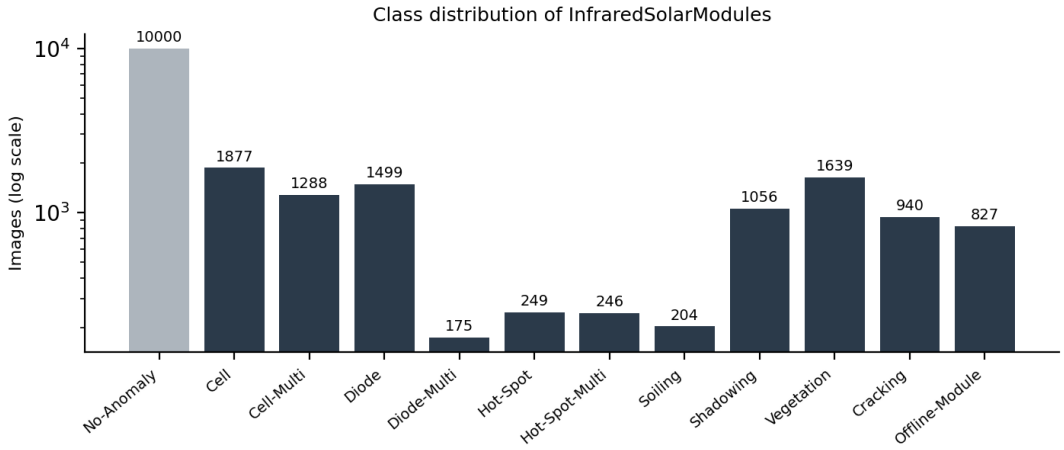


Fig. 2. Class distribution of the InfraredSolarModules dataset (logarithmic axis).

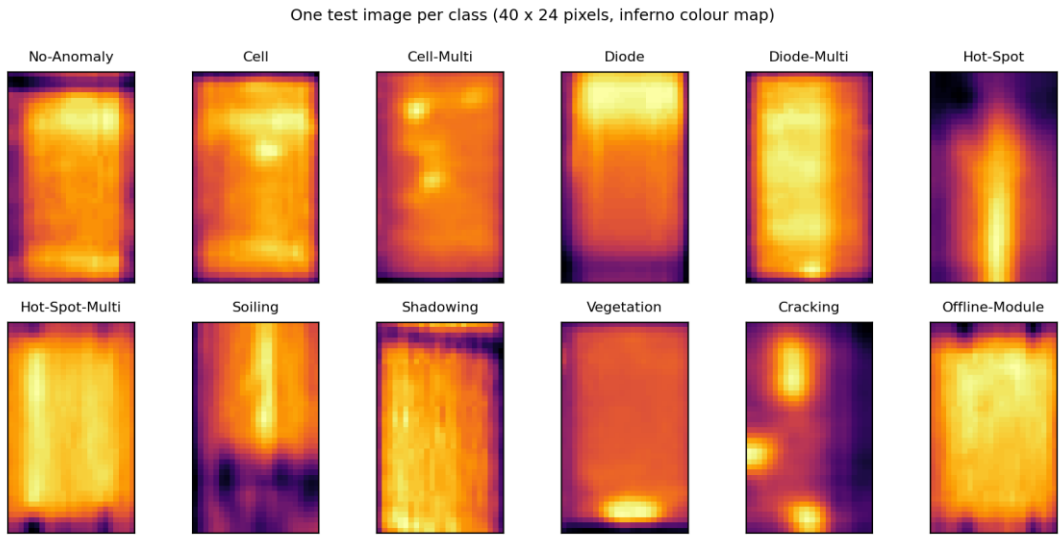


Fig. 3. One test image per class, shown with the inferno colour map.

5.2.2 Split and Normalisation

The images are split once into 70% training (14,000), 15% validation (3,000) and 15% test (3,000) images, stratified by class with a fixed random seed, so that every model sees exactly the same split (Table 3.1). Pixel values are scaled to [0, 1] and normalised with statistics computed on the training split only, which prevents information from the validation and test images from entering the pipeline:

$$x' = (x - \mu_{train}) / \sigma_{train}, \quad \mu_{train} = 0.618, \quad \sigma_{train} = 0.156$$

5.2.3 Augmentation

During training only random horizontal and vertical flips (probability 0.5 each) are applied. A flip does not change the class of a module, because the faults are defined by which cells or diodes are hot and not by orientation. Brightness, contrast and noise jitter were implemented and tested but removed, because in thermal images brightness is temperature, and perturbing it lowered accuracy by about ten points (Section 6.3.1). Quarter-turn rotations are not used, because they would change the 40 x 24 shape and the convolutional pipeline assumes a fixed size.

5.2.4 Handling Class Imbalance

Class imbalance is handled at three levels: structurally by the early exit of Stage 1, in the loss by class weights, and in the evaluation by macro-averaged metrics. For a network with K

classes trained on N samples, class r with n_r samples receives the weight

$$w_r = N / (K \times n_r)$$

so that every class contributes equally to the loss in expectation. For example, in the flat 12-class setting the healthy class receives a weight of 0.17 and Diode-Multi a weight of 9.5.

5.3 Building Blocks and Their Mathematics

All networks are assembled from a small set of layers. This section states the operation of each layer, so that the architecture tables and parameter counts in the following sections can be checked by hand.

5.3.1 Convolution

A convolutional layer with C_in input channels, C_out output channels and a k x k kernel computes, for every output channel c and position (i, j),

$$y[c, i, j] = b[c] + \text{sum over } (m, u, v) \text{ of } W[c, m, u, v] * x[m, i + u, j + v]$$

with padding that preserves the spatial size (k = 3, padding 1). The layer has k^2 x C_in x C_out weights and C_out biases (the bias is omitted when a BatchNorm layer follows).

5.3.2 Batch Normalisation

BatchNorm normalises each channel over the mini-batch and then rescales it with two learned parameters per channel:

$$x_hat = (x - mean_B) / sqrt(var_B + eps), \quad y = gamma * x_hat + beta$$

It stabilises and speeds up training and reduces sensitivity to weight initialisation, which matters for small networks trained on imbalanced data. At test time running averages of the mean and variance are used.

5.3.3 ReLU and Max Pooling

The rectified linear unit f(x) = max(0, x) introduces the non-linearity that lets the network learn complex decision boundaries. Max pooling with a 2 x 2 window and stride 2 keeps the maximum of each window, halving height and width, reducing computation and giving tolerance to small shifts in where the heat pattern sits in the image.

5.3.4 Global Average Pooling and Classification Head

After the last block each of the C feature maps of size H x W is summarised by its mean, z_c = (1 / HW) sum of x_c(i, j). Dropout with probability p randomly zeroes activations during training as a regulariser, and a fully connected layer maps the C values to K class scores s_k. The class probabilities are given by the softmax function:

$$p_k = exp(s_k) / sum\ over\ j\ of\ exp(s_j)$$

5.3.5 Loss Function

Each network minimises the class-weighted cross-entropy over a mini-batch of B samples with labels y_i:

$$L = - (1 / B) * sum\ over\ i\ of\ w[y_i] * log\ p[i, y_i]$$

5.3.6 Counting Parameters

A block of two convolutions followed by BatchNorm has k^2 x C_in x C_mid + k^2 x C_mid x C_out weights and 2 x (C_mid + C_out) BatchNorm parameters. The tables in the following sections were generated by running the implemented models, and the totals match this formula.

5.4 Stage 1: CNN-1, Binary Triage

CNN-1 answers whether a module is healthy or defective. In the from-scratch variant it has three blocks of two 3 x 3 convolutions, each followed by BatchNorm and ReLU, then 2 x 2 max pooling, with 64, 128 and 128 channels, followed by global average pooling, dropout (0.2) and a fully connected layer with two outputs, 555,074 parameters in total. The 80 x 48 upsampled input is reduced to 64 x 40 x 24, 128 x 20 x 12 and 128 x 10 x 6 feature maps. A prediction of No-Anomaly ends the pipeline (the early exit); a defect continues to Stage 2. Table 5.1 lists every layer.

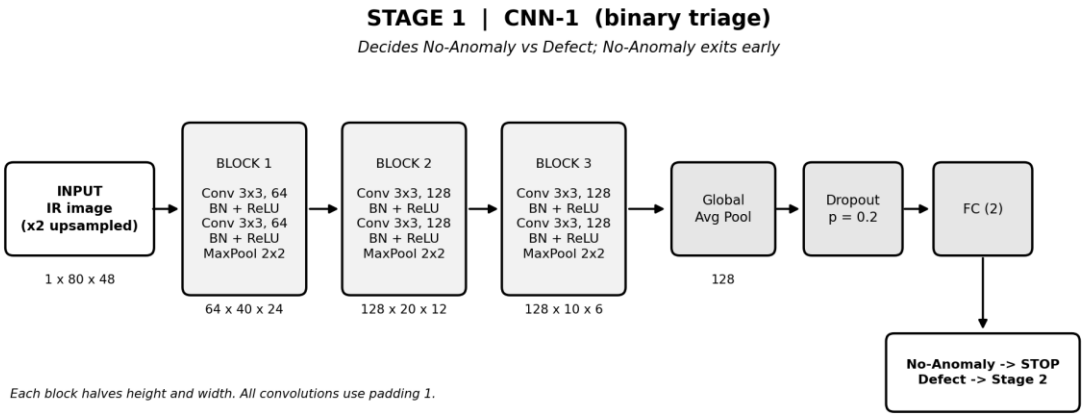


Fig. 4. Stage 1: CNN-1, binary triage (layers and feature-map sizes).

Table 5.1: Layer-by-layer structure of CNN-1 (from-scratch variant)

Layer	Output shape	Parameters
Input	1 x 80 x 48	0
Conv2d(64)	64 x 80 x 48	576
BatchNorm2d	64 x 80 x 48	128
ReLU	64 x 80 x 48	0
Conv2d(64)	64 x 80 x 48	36,864
BatchNorm2d	64 x 80 x 48	128
ReLU	64 x 80 x 48	0
MaxPool2d	64 x 40 x 24	0
Conv2d(128)	128 x 40 x 24	73,728
BatchNorm2d	128 x 40 x 24	256
ReLU	128 x 40 x 24	0
Conv2d(128)	128 x 40 x 24	147,456
BatchNorm2d	128 x 40 x 24	256
ReLU	128 x 40 x 24	0
MaxPool2d	128 x 20 x 12	0
Conv2d(128)	128 x 20 x 12	147,456
BatchNorm2d	128 x 20 x 12	256
ReLU	128 x 20 x 12	0
Conv2d(128)	128 x 20 x 12	147,456
BatchNorm2d	128 x 20 x 12	256
ReLU	128 x 20 x 12	0
MaxPool2d	128 x 10 x 6	0
GlobalAvgPool	128	0
FC(2)	2	258
Total		**555,074**

5.5 Stage 2: CNN-2, Fault-Family Classifier

CNN-2 has the same three-block design with 128, 256 and 128 channels, dropout 0.4 and three outputs, 1,478,147 parameters in total. It is trained only on defective images labelled by family. Its highest-scoring output selects which specialist runs next. Faults of one family should look

alike, so the decision should be easier than choosing among eleven faults directly. Section 6.4 shows how well this holds.

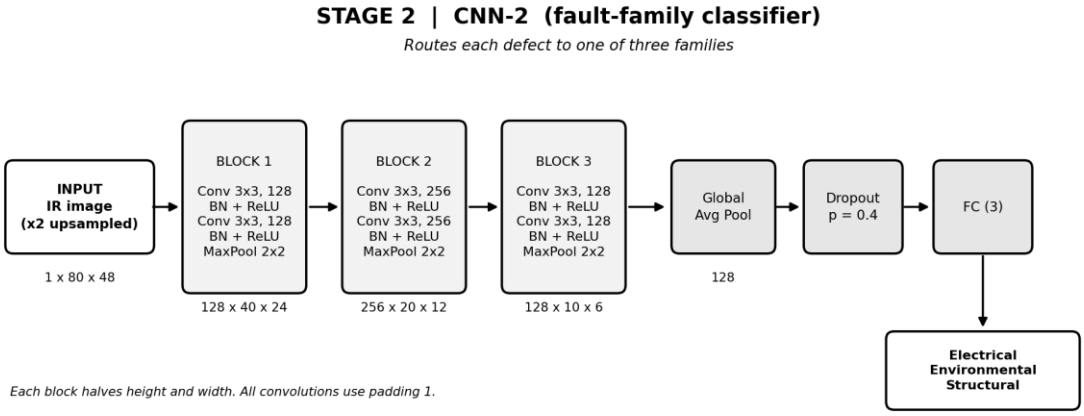


Fig. 5. Stage 2: CNN-2, fault-family classifier.

Table 5.2: Layer-by-layer structure of CNN-2 (from-scratch variant)

Layer	Output shape	Parameters
Input	1 x 80 x 48	0
Conv2d(128)	128 x 80 x 48	1,152
BatchNorm2d	128 x 80 x 48	256
ReLU	128 x 80 x 48	0
Conv2d(128)	128 x 80 x 48	147,456
BatchNorm2d	128 x 80 x 48	256
ReLU	128 x 80 x 48	0
MaxPool2d	128 x 40 x 24	0
Conv2d(256)	256 x 40 x 24	294,912
BatchNorm2d	256 x 40 x 24	512
ReLU	256 x 40 x 24	0
Conv2d(256)	256 x 40 x 24	589,824
BatchNorm2d	256 x 40 x 24	512
ReLU	256 x 40 x 24	0
MaxPool2d	256 x 20 x 12	0
Conv2d(128)	128 x 20 x 12	294,912
BatchNorm2d	128 x 20 x 12	256
ReLU	128 x 20 x 12	0
Conv2d(128)	128 x 20 x 12	147,456
BatchNorm2d	128 x 20 x 12	256
ReLU	128 x 20 x 12	0
MaxPool2d	128 x 10 x 6	0
GlobalAvgPool	128	0
FC(3)	3	387
Total		**1,478,147**

5.6 Stage 3: Fine-Grained Specialists

Three specialists share the 128-256-128 design and differ only in their output layer: CNN-3a

(Electrical, six classes), CNN-3b (Environmental, three classes) and CNN-3c (Structural, two classes). Each is trained only on the defective images of its own family, so it can spend all its capacity on the distinctions inside that family. The final label is the family's class offset plus the specialist's prediction. Table 5.3 gives the parameter counts of all five networks in both variants.

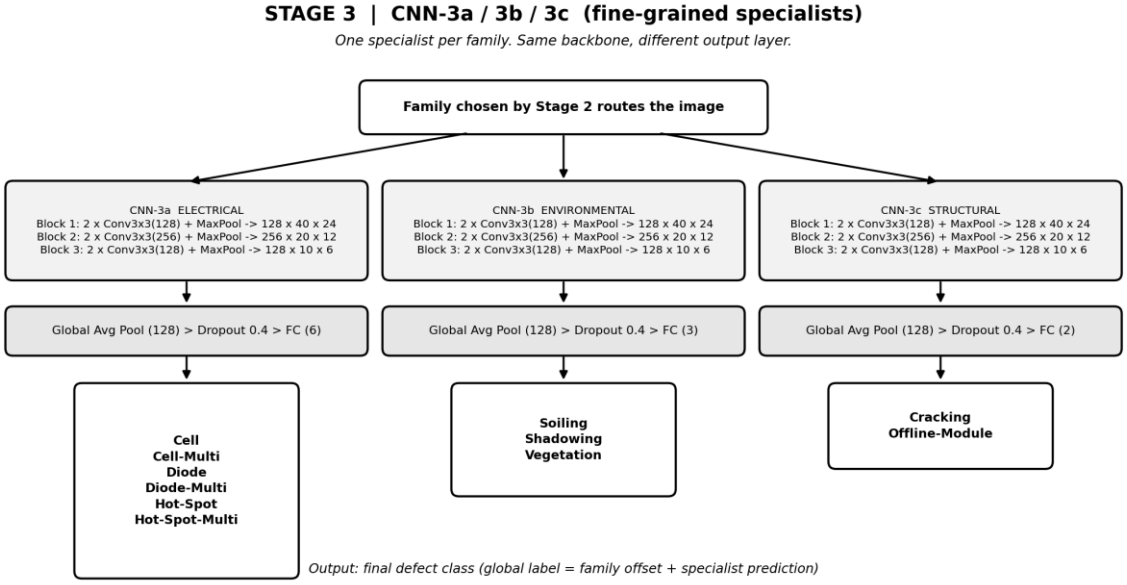


Fig. 6. Stage 3: the three specialists and their classes.

Table 5.3: Parameter counts of the five stage networks

Network	Output classes	From scratch	Pretrained EfficientNet-B0
CNN-1 (triage)	2	555,074	4,010,110
CNN-2 (family)	3	1,478,147	4,011,391
CNN-3a (Electrical)	6	1,478,534	4,015,234
CNN-3b (Environmental)	3	1,478,147	4,011,391
CNN-3c (Structural)	2	1,478,018	4,010,110
Cascade total		**6,467,920**	**20,058,236**

Table 5.4: Architecture summary of the five networks (from-scratch variant)

Network	Stage	Input	Layers	Output
CNN-1	1: triage	IR image, 80 x 48	3 blocks 2 x (Conv3x3, BN, ReLU) + MaxPool, 64-128-128; GAP, Dropout 0.2, FC	2 classes
CNN-2	2: family	Defective image	Same, 128-256-128; Dropout 0.4, FC	3 families
CNN-3a	3: Electrical	Routed image	Same as CNN-2, FC with 6 outputs	6 classes
CNN-3b	3: Environmental	Routed image	Same as CNN-2, FC with 3 outputs	3 classes
CNN-3c	3: Structural	Routed image	Same as CNN-2, FC with 2 outputs	2 classes

5.7 Pretrained Variant

In the pretrained variant each of the five networks is an ImageNet-pretrained EfficientNet-B0 [11] whose classifier is replaced by dropout and a fully connected layer with the stage's number of outputs. EfficientNet-B0 is built from mobile inverted bottleneck (MBConv) blocks: each block expands the channels with a 1 x 1 convolution, applies a depthwise convolution, re-weights channels with a squeeze-and-excitation module and projects back with another 1 x 1 convolution, adding a residual connection when the shape allows. The family of networks was designed by compound scaling of depth, width and resolution, which gives B0 about 4 million parameters in the form used here.

The grayscale image is resized to 160 x 96 with bilinear interpolation and repeated to three channels inside the model, so the rest of the pipeline is unchanged and the pretrained first layer can be used as is. The five networks hold 20.1 million parameters in total, against 6.5 million for the from-scratch cascade. All weights are fine-tuned, with a learning rate of 3e-4 and mixed-precision arithmetic.

5.8 Training Procedure

Each network is trained independently on the data matching its task (ground-truth routing). Table 5.5 shows the training and validation sizes of the five tasks. Training uses the Adam optimiser with weight decay 1e-4, batch size 64 and the class-weighted loss of Section 5.3.5. The learning rate is 5e-4 for the from-scratch networks and 3e-4 for the pretrained ones. The checkpoint with the best validation macro-F1 is kept, the learning rate is halved when that metric stalls, and training stops after 25 epochs (6 for the pretrained networks) without improvement.

Table 5.5: Training data of each stage network

Network	Task	Training images	Validation images	Classes
CNN-1	Healthy vs defect	14,000	3,000	2
CNN-2	Family of a defect	7,000	1,500	3
CNN-3a	Electrical subtype	3,734	801	6
CNN-3b	Environmental subtype	2,029	434	3
CNN-3c	Structural subtype	1,237	265	2

Algorithm 5.1: Training of one stage network

```
input : training set D_train, validation set D_val, network f, class weights w
output: best checkpoint f*
best <- -1;  wait <- 0
for epoch = 1 .. E_max:
```

```
    for each mini-batch (x, y) of D_train:
        x <- normalise(random_flip(x))
        loss <- weighted_cross_entropy(f(x), y, w)
        update the parameters of f with Adam on loss
    F1_val <- macro_F1(f on D_val)
    if F1_val > best:  best <- F1_val;  save checkpoint f*;  wait <- 0
    else:              wait <- wait + 1;  if wait >= patience: stop
return f*
```

5.9 Cascade Inference and Cost Analysis

Algorithm 5.2: Inference of Cascaded-SolarNet on one image

```
input : image x
output: class label y
x <- normalise(x)
if CNN-1(x) = No-Anomaly:  return No-Anomaly          # early exit
family <- argmax CNN-2(x)
k      <- argmax CNN-3[family](x)
return class_offset[family] + k
```

Let c1, c2 and c3 be the costs of the three kinds of network. A healthy image costs c1 and a defective one c1 + c2 + c3, so the expected cost per image is c1 + P(defect)(c2 + c3), where P(defect) is the fraction of images that pass Stage 1. On the test split 50.8% of the images leave at Stage 1. With similar-sized networks in all stages the cascade therefore costs about 1.5 network evaluations per image on average, instead of one for a flat model, and it only saves time when the first stage is much cheaper than the others. In the from-scratch variant CNN-1 has about 40% of the parameters of the others, but in the pretrained variant all five networks are the same size, which is why the early exit gave no efficiency advantage in our measurements.

5.10 Evaluation Tools and Demonstrations

The evaluation script runs the whole cascade end to end on the test split and reports accuracy, macro-F1, Cohen's kappa, per-class scores, the confusion matrix and the error attribution per stage (missed defects and false alarms at Stage 1, wrong family at Stage 2, wrong subtype at Stage 3). Further scripts reproduce the baselines, run 10-fold cross-validation, search for the data-driven family grouping and generate the figures and reports. Two demonstrations show the cascade at work: a command-line tool that draws each stage's confidence bars for chosen test images (Figure 7), and an interactive browser page in which a class can be selected and the three stages reveal their decisions one after another. Appendix C explains how to run them.

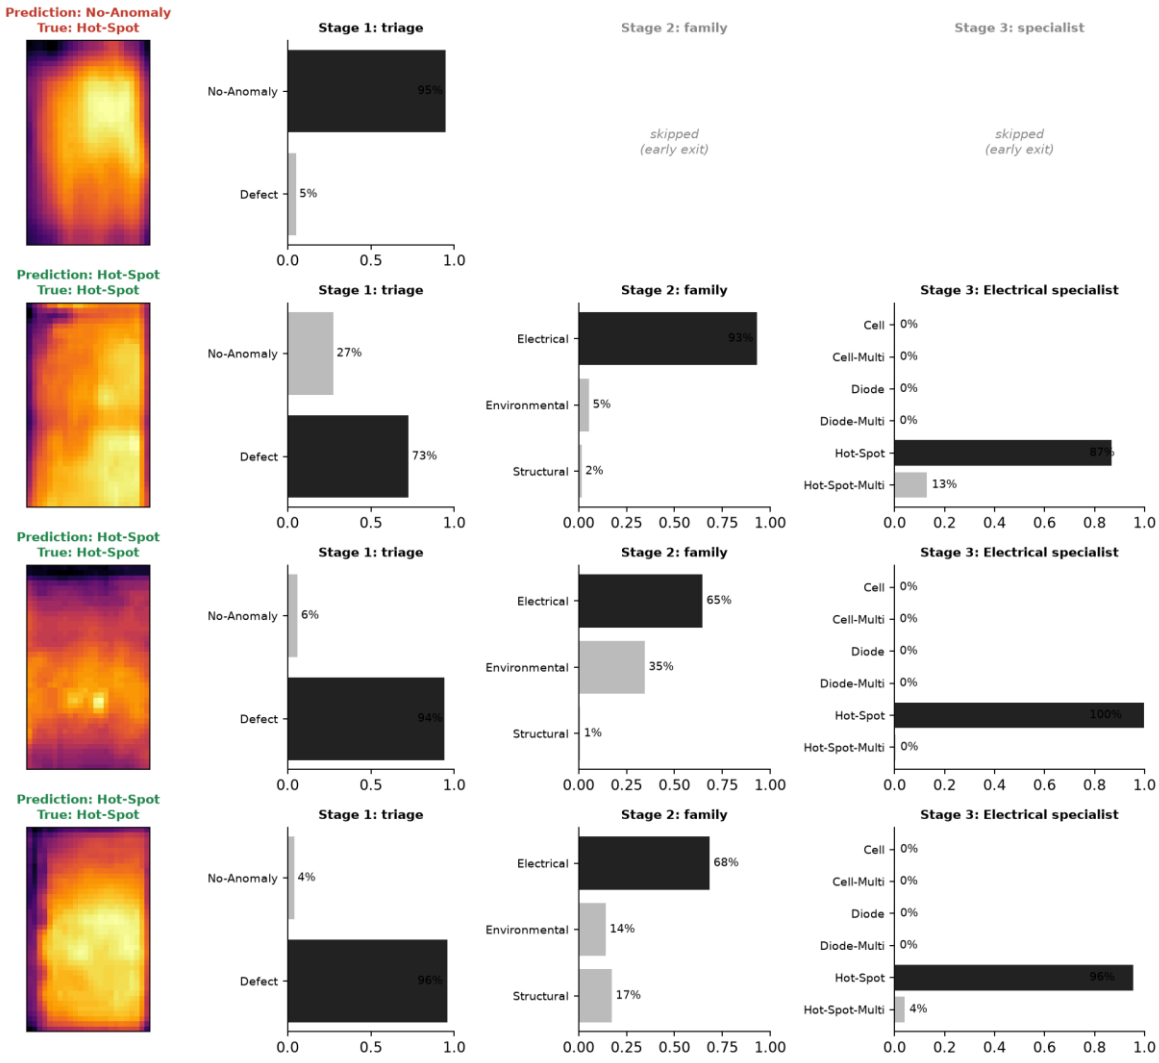


Fig. 7. Output of the live demonstration: each stage's decision and confidence for Hot-Spot test images.

5.11 Testing and Reproducibility

A suite of 15 automated tests runs on synthetic data and needs no dataset download. It verifies the class hierarchy, the data split, the augmentation, the shapes of all networks and the routing of the cascade. Table 5.6 lists the tests. Reproducibility is supported by a fixed seed, a saved split file, configuration files for every experiment and result files for every reported number.

Table 5.6: Automated tests

Test	What it verifies
Twelve classes, healthy first	The taxonomy has 12 classes and the healthy class has global index 0
Node labels cover every class once	Each class belongs to exactly one family and receives a consistent local label

Test	What it verifies
Global index matches family offsets	Family offset plus local index reproduces the global class index
Duplicate class rejected	A taxonomy listing a class twice raises an error
Class weights formula	Weights equal N / (K x n) and sum correctly
Stratified split	Train, validation and test ids are disjoint, complete and reproducible for a fixed seed
Node output shape (5 nodes)	Each stage network outputs one score per class for a 40 x 24 input
Augmentation shape and range	Flips and jitter keep the tensor shape and values within [0, 1]
Cascade routing and early exit	Healthy images exit at Stage 1 and defects reach the correct specialist and offset
Pretrained stage output shape	The EfficientNet-B0 stage outputs the right number of classes

5.12 Chapter Summary

The system consists of five small stage networks in two variants, a leak-free data pipeline, a class-weighted training procedure, an evaluation module that attributes every error to a stage and an automated test suite. The next chapter evaluates it.

CHAPTER 6

RESULTS AND ANALYSIS

6.1 Experimental Setup

6.1.1 Dataset and Protocol

All numbers are measured on the 3,000-image held-out test split of Table 3.1 and are single runs per configuration (seed 42). The test images are not used for training, early stopping, normalisation or model selection. Choices such as the learning rate were made on the validation split.

6.1.2 Implementation and Training Configuration

Table 6.1: Training configuration of the evaluated models

Group	Parameter	Value
Data	Split	Stratified 70/15/15 (14,000 / 3,000 / 3,000), seed 42
Data	Normalisation	Mean and standard deviation of the training split
Data	Augmentation	Random horizontal and vertical flips (p = 0.5)
From scratch	Input size	80 x 48 (2x bilinear upsampling)
From scratch	Blocks and channels	3 blocks of 2 x (Conv3x3, BN, ReLU) + MaxPool; 64-128-128 (Stage 1), 128-256-128 (others)
From scratch	Dropout	0.2 (Stage 1), 0.4 (Stage 2 and 3)
From scratch	Learning rate, patience	5e-4; 25 epochs (maximum 150)
Pretrained	Backbone	EfficientNet-B0, ImageNet weights, new head
Pretrained	Input size	160 x 96, channel repeated to 3
Pretrained	Dropout	0.2 (Stage 1), 0.3 (others)
Pretrained	Learning rate, patience	3e-4; 6 epochs (maximum 25), mixed precision
Common	Optimiser	Adam, weight decay 1e-4, batch size 64
Common	Loss	Class-weighted cross-entropy, w = N / (K x n)
Common	Model selection	Early stopping on validation macro-F1
Common	Hardware, software	NVIDIA RTX 5060 Laptop GPU; PyTorch 2.11

6.1.3 Evaluation Metrics

With TP, FP and FN the true positives, false positives and false negatives of a class, precision is TP / (TP + FP), recall is TP / (TP + FN) and the F1 score is their harmonic mean:

$$F1 = 2 \times precision \times recall / (precision + recall), \quad macro\text{-}F1 = (1 / K) \times \sum_{\text{classes}} F1_r$$

Macro-F1 weights every class equally and therefore exposes poor performance on rare faults, which accuracy hides under class imbalance. Cohen's kappa measures agreement beyond

chance:

$$kappa = (p_o - p_e) / (1 - p_e)$$

where p_o is the observed accuracy and p_e the accuracy expected by chance from the class frequencies. The Stage 1 binary accuracy measures the healthy-versus-defect decision, and per-stage error counts attribute each mistake to the first stage that went wrong.

6.1.4 Configurations

Table 6.2: Definition of the evaluated configurations

ID	Variant	Modification
C1	Default	Flips plus brightness (20%), contrast (15%) and noise jitter; 40 x 24 input; 60 epochs
C2	No photometric jitter	C1 with flips only; longer schedule (100 epochs)
C3	Improved	C2 with 80 x 48 input and two convolutions per block; 150 epochs
C4	Improved2	C3 with balanced sampling and label smoothing instead of loss weights
C5	Improved3	C3 with larger Stage 1 and Stage 2 networks
C6	Data-driven families	C5 with the data-driven family grouping of Section 4.3
C7	Pretrained	EfficientNet-B0 (ImageNet) fine-tuned at every stage

6.2 Overall Performance

6.2.1 Test Results

The best result is the pretrained cascade (C7) with 83.2% accuracy, macro-F1 0.700, Cohen's kappa 0.764 and 95.8% on the Stage 1 healthy-versus-defect decision. The best from-scratch cascade reaches 79.8% accuracy. Table 6.3 lists all configurations and baselines, and Figure 8 plots them.

Table 6.3: Test results by configuration

Configuration	Model	Accuracy	Macro-F1	Kappa	Stage 1 binary
C1 Default (jitter on)	Cascade	61.7%	0.368	0.461	81.0%
	Flat CNN	47.6%	0.379	0.357	n/a
C2 No photometric jitter	Cascade	71.1%	0.509	0.583	86.0%
	Flat CNN	68.5%	0.571	0.586	n/a
C3 Improved (80x48, deeper)	Cascade	76.0%	0.582	0.656	89.5%
	Flat CNN	76.8%	0.653	0.689	n/a
C4 Improved2 (balanced sampling)	Cascade	75.5%	0.576	0.648	89.1%
C5 Improved3 (bigger Stage 1/2)	Cascade	79.6%	0.638	0.714	93.8%
C6 Data-driven families	Cascade	79.8%	0.646	0.717	94.1%
C7 Pretrained EfficientNet-B0	**Cascade**	**83.2%**	**0.700**	**0.764**	**95.8%**
	Flat CNN	82.6%	0.708	0.761	n/a
Raw-pixel Random Forest	RF	61.6%	0.337	0.324	n/a
EfficientNet-B0 features + SVM	SVM	78.7%	0.640	0.689	90.4%

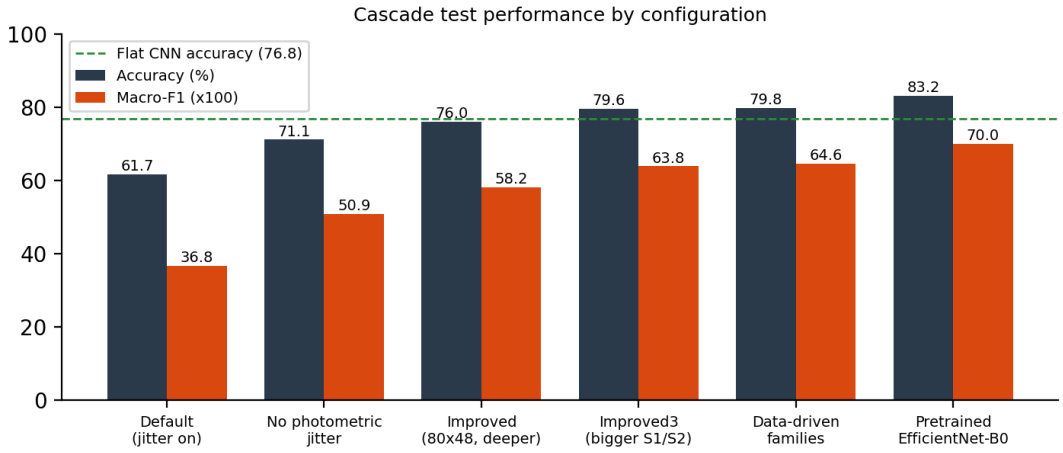


Fig. 8. Cascade test accuracy and macro-F1 by configuration (dashed line: from-scratch flat CNN).

6.2.2 Per-Network Validation Results and Training Behaviour

Figure 9 compares the best validation macro-F1 of each of the five networks in the two variants, and Figure 10 shows the training curves. Pretraining improved every network, most clearly Stage 2 (validation macro-F1 0.717 to 0.796) and the Electrical specialist (0.786 to 0.826). The pretrained networks reach their best validation score within 10 to 20 epochs, while the from-scratch networks need 50 to 140 epochs, and the validation curves of the from-scratch family and electrical networks are noticeably noisier.

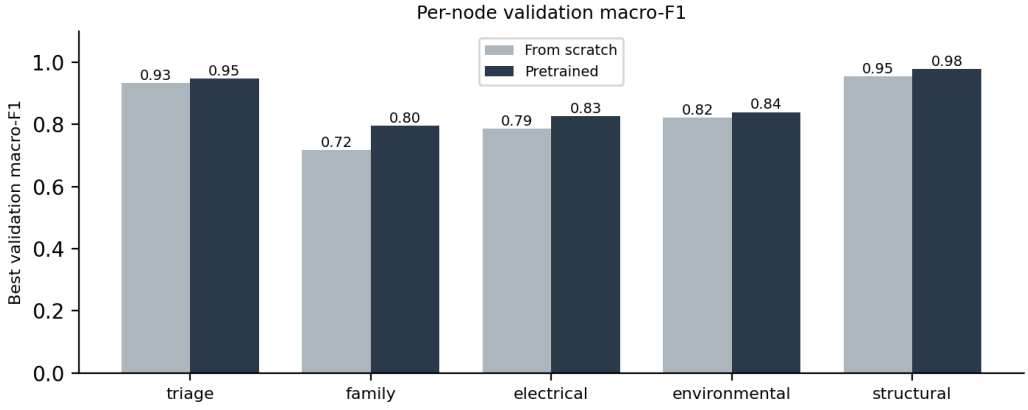


Fig. 9. Best validation macro-F1 of each stage network, from scratch and pretrained.

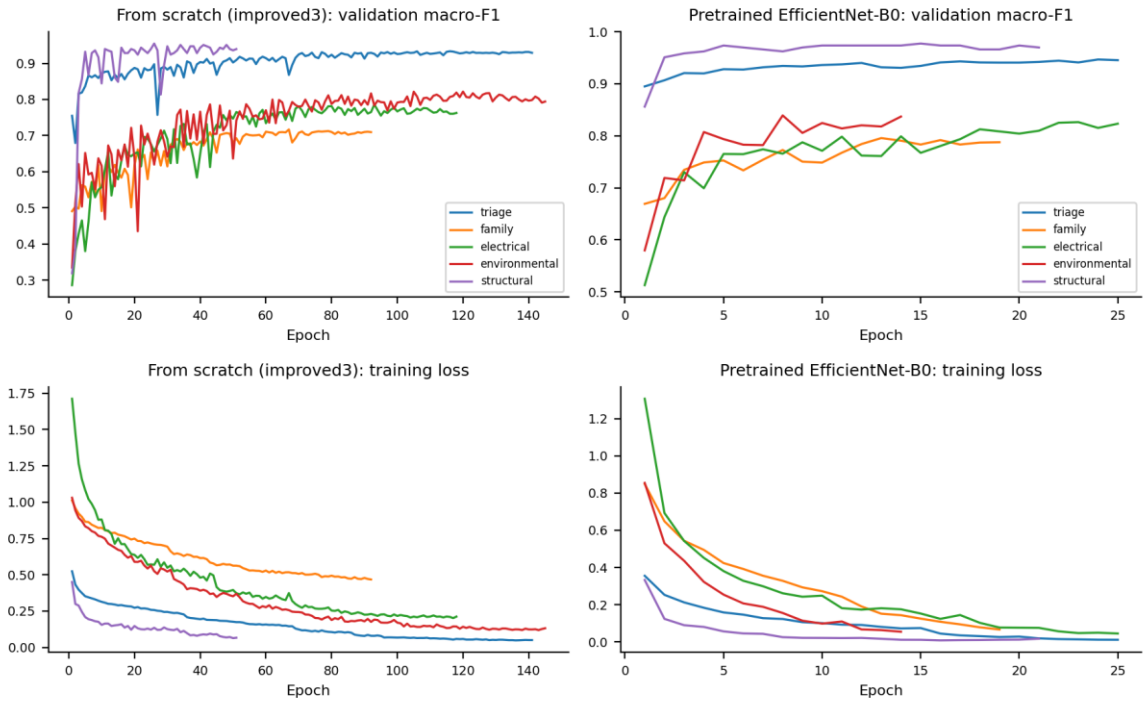


Fig. 10. Validation macro-F1 and training loss of the five stage networks.

6.3 Ablation Study

The ablations change one design choice at a time, in the order in which the project progressed (Table 6.2). Each subsection reports what changed and what was measured.

6.3.1 Photometric Augmentation

Table 6.4: Effect of photometric augmentation (cascade, 40 x 24 input)

Configuration	Accuracy	Macro-F1	Kappa	Stage 1 binary
C1 with brightness, contrast and noise jitter	61.7%	0.368	0.461	81.0%
C2 flips only	71.1%	0.509	0.583	86.0%

Removing the jitter raised accuracy by about ten points and cut Stage 1 false alarms from 250 to 118. In an infrared image the absolute brightness of a region is its temperature, which is exactly what separates a hot spot from a healthy module, so scaling or shifting brightness changes the evidence and not only the appearance. This was the largest single gain of the project and the main lesson about thermal data.

6.3.2 Input Size and Network Capacity

Table 6.5: Effect of input size and capacity (cascade, flips only)

Configuration	Accuracy	Macro-F1	Kappa	Stage 1 binary
C2 40 x 24 input	71.1%	0.509	0.583	86.0%
C3 80 x 48 input, two convolutions per block	76.0%	0.582	0.656	89.5%
C5 plus larger Stage 1 and Stage 2 networks	79.6%	0.638	0.714	93.8%

Upsampling the image to 80 x 48 and adding a second convolution per block gave 76.0% accuracy, and enlarging the two front stages gave 79.6%. Stage 1 missed defects fell from 233 to 108. These changes target the early stages because their errors cannot be corrected later.

6.3.3 Balanced Sampling and Label Smoothing

Replacing the class-weighted loss by sampling with weight proportional to 1/sqrt(class count) and adding label smoothing of 0.05 (C4) gave 75.5% accuracy and macro-F1 0.576, slightly below the weighted loss (76.0%, 0.582). The class-weighted loss was therefore kept.

6.3.4 Inference-Time Variants

Table 6.6: Inference-time variants on the improved cascade (no retraining)

Variant	Accuracy	Macro-F1	Kappa
Hard routing (baseline)	76.0%	0.582	0.656
Flip test-time augmentation	76.4%	0.582	0.660
Stage 1 threshold tuned on validation data	76.0%	0.582	0.656
Soft routing (multiplied probabilities)	75.9%	0.575	0.648
Soft routing with test-time augmentation	76.3%	0.579	0.653

None of the variants helped meaningfully. The validation-optimal Stage 1 threshold stayed at 0.5, so tuning it changed nothing. Soft routing, which multiplies the probabilities of the stages so that a doubtful family decision can be outvoted, did not recover the Stage 2 errors, which suggests that those errors are confident mistakes and not close calls. Test-time augmentation added 0.3 points at four times the cost.

6.3.5 Fault-Family Grouping

Table 6.7 compares the semantic grouping with the data-driven grouping of Section 4.3. The data-driven grouping reduced wrong-family errors at Stage 2 from 300 to 116, but moved the difficulty into the fine-grained stage, where its largest family now contains the visually similar classes Cell, Cell-Multi, Soiling, Vegetation and Cracking. The net effect on accuracy was

small.

Table 6.7: Semantic versus data-driven family grouping (from-scratch cascade)

Grouping	Stage 2 errors	Stage 3 errors	Total errors	Accuracy	Macro-F1
Semantic (Electrical, Environmental, Structural)	300	126	613	79.6%	0.638
Data-driven (groups A, B, C)	116	313	605	79.8%	0.646

6.3.6 ImageNet Pretraining

Fine-tuning ImageNet-pretrained EfficientNet-B0 at every stage raised the cascade from 79.8% to 83.2% accuracy and macro-F1 from 0.646 to 0.700, and training took minutes instead of an hour. The pretrained flat CNN improved even more, from 76.8% to 82.6%, so that the two designs end up tied. Pretraining helps on these small thermal images because the early filters (edges, blobs, gradients) transfer, which refutes the expectation in the project proposal that pretrained backbones would give little benefit.

6.3.7 Summary of the Ablations

In order of effect: removing photometric jitter (+9.4 points), pretraining (+3.4 and +5.8 for the flat model), larger front stages and input (+3.6 and +5.0 on the way), while sampling changes, inference variants and test-time augmentation gave nothing. The changes that helped were those that preserved or exposed the thermal signal or brought in better features. The changes that did not help tried to repair Stage 2 after the fact.

6.4 Error Analysis by Stage

Table 6.8 attributes each test error to the first stage that went wrong.

Table 6.8: Test errors per stage (out of 3,000 images)

Configuration	Stage 1 missed defect	Stage 1 false alarm	Stage 2 wrong family	Stage 3 wrong subtype	Total
C3 Improved	233	83	294	109	719
C5 Improved3	108	79	300	126	613
C6 Data-driven	110	66	116	313	605
C7 Pretrained	**75**	**52**	**272**	**106**	**505**

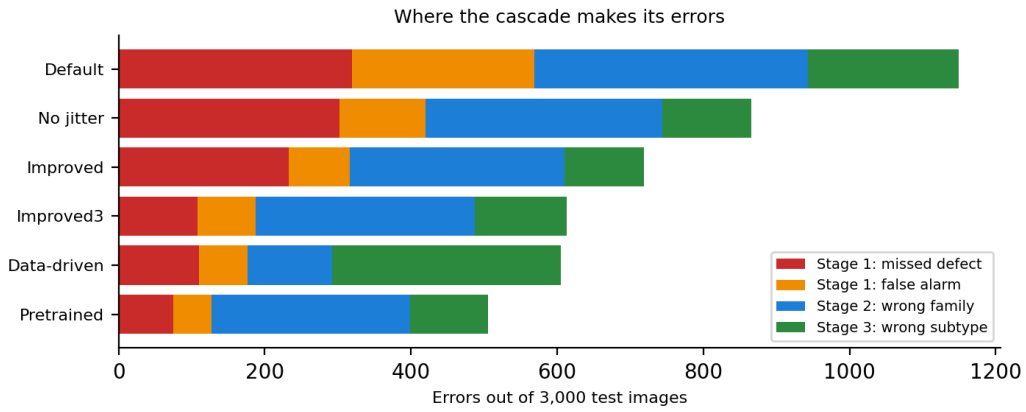


Fig. 11. Where the cascade makes its errors, by configuration.

Pretraining reduced errors at every stage compared with C5 (total 613 to 505), but Stage 2, the family decision, remains the largest source (272 of 505 errors), and an error there cannot be corrected by Stage 3. About 50.8% of the test images leave the cascade at Stage 1. Stage 1 itself is accurate: only 75 defects were missed and 52 healthy modules flagged, a binary accuracy of 95.8%. The missed defects matter most in practice, because a missed hot spot is a safety risk, whereas a false alarm costs only an inspection.

6.5 Per-Class Performance and Confusion

6.5.1 Per-Class F1

The rare classes are the weak point. They have only 123 to 174 training images, and several look alike. Class weighting, stratified splitting and macro-F1 reduce but do not remove the problem.

Table 6.9: Per-class F1 on the test set

Class	Test images	Cascade, from scratch (C5)	Cascade, pretrained (C7)	Flat CNN, pretrained
No-Anomaly	1500	0.94	0.96	0.94
Cell	281	0.63	0.69	0.69
Cell-Multi	193	0.47	0.52	0.49
Diode	225	0.93	0.96	0.96
Diode-Multi	26	0.82	0.87	0.91
Hot-Spot	37	0.54	0.65	0.62
Hot-Spot-Multi	37	0.57	0.65	0.62
Soiling	31	0.22	0.30	0.35
Shadowing	159	0.67	0.72	0.74
Vegetation	246	0.65	0.65	0.71
Cracking	141	0.65	0.70	0.75
Offline-Module	124	0.59	0.74	0.71

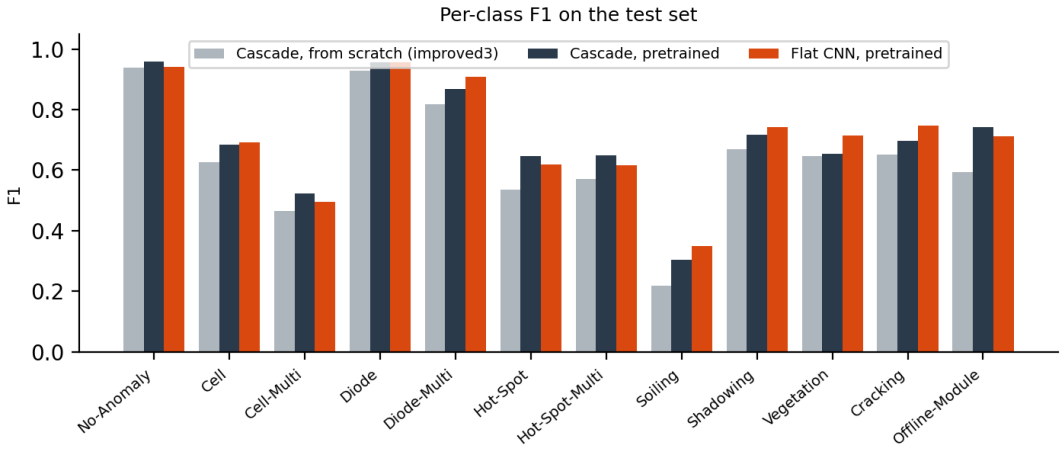


Fig. 12. Per-class F1 on the test set.

6.5.2 Confusion Matrices

Figures 13 and 14 show the row-normalised confusion matrices of the pretrained cascade and the pretrained flat CNN. Both are diagonal for the healthy class and for Diode, Diode-Multi and Cracking, and both spread the probability mass of Cell, Cell-Multi, Hot-Spot and Soiling over several neighbouring classes. Table 6.10 lists the most frequent confusions of the cascade.

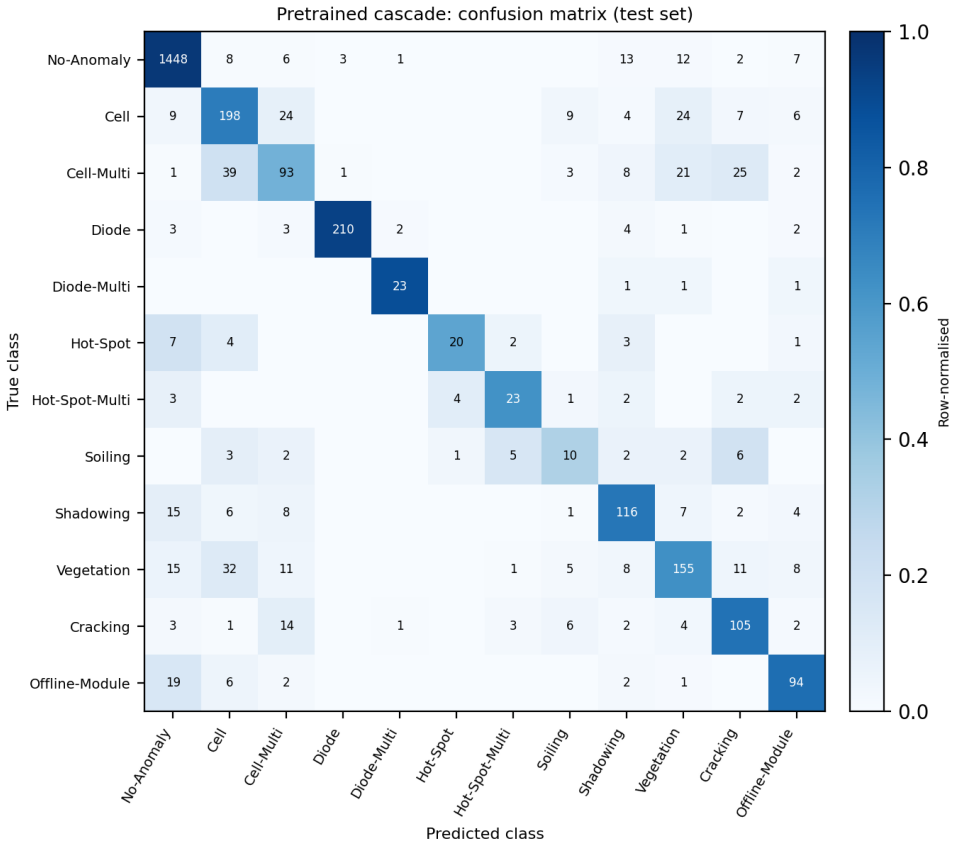


Fig. 13. Confusion matrix of the pretrained cascade on the test set (row-normalised, counts inside cells).

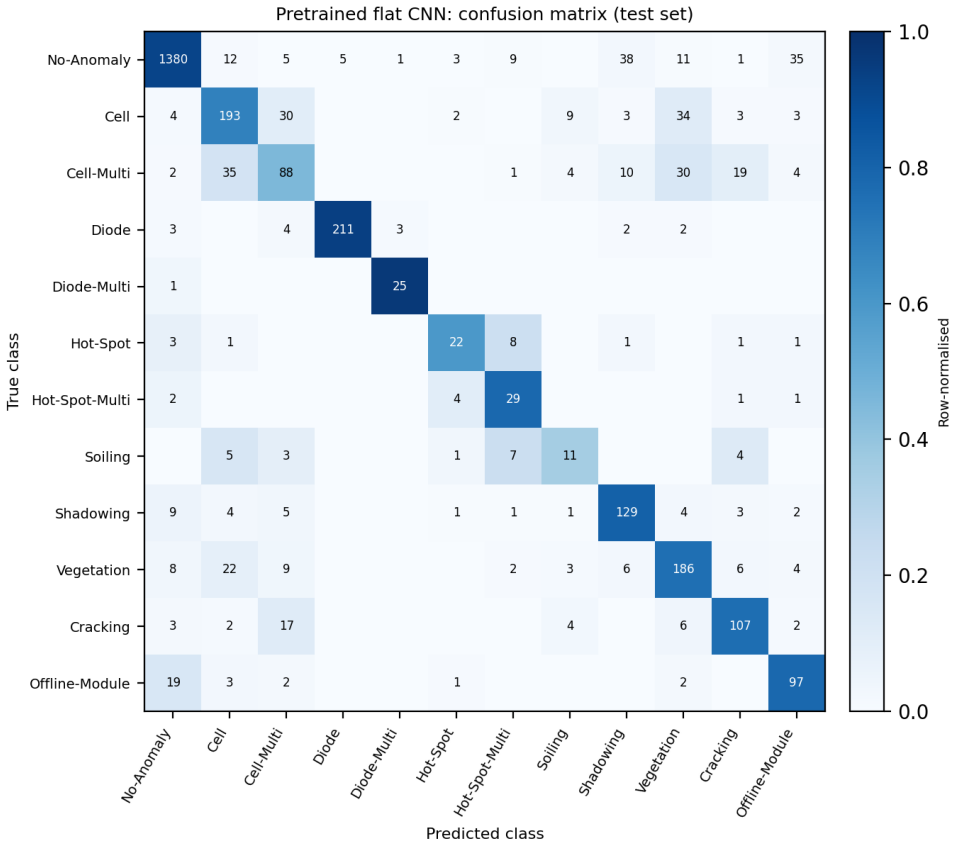


Fig. 14. Confusion matrix of the pretrained flat CNN on the test set.

Table 6.10: Most frequent confusions of the pretrained cascade

True class	Predicted as	Images	Share of the true class
Cell-Multi	Cell	39	20%
Vegetation	Cell	32	13%
Cell-Multi	Cracking	25	13%
Cell	Vegetation	24	9%
Cell	Cell-Multi	24	9%
Cell-Multi	Vegetation	21	11%

The most frequent confusions are between classes of the same physical phenomenon (a single versus several affected cells or diodes) and between faults and the healthy class, in particular for faint hot spots. These are also the cases where a human inspector would hesitate in a 24 x 40 pixel image.

6.6 Comparison with Published Work

The dataset paper reports no model accuracies, and the figures often attributed to it (Random Forest 84%, CNN 92.5%, EfficientNet + SVM 93%) come from later papers that differ in task,

protocol and features. Table 6.11 therefore places published results next to ours but does not rank them.

Table 6.11: Published results on InfraredSolarModules and results of this project

Paper	Method	Task and protocol	Reported	Checked
Duranay 2023 [2]	EfficientNet-B0 features (ImageNet), NCA selection, SVM	12 classes, 10-fold cross-validation	93.93% accuracy, 89.82% F1	Full paper read
Alves et al. 2021 [3]	Custom CNN with augmentation	Binary detection; 8 selected classes; cross-validation	92.5% binary; 78.85% for 8 classes	Abstract only
Corley et al. 2025 [4]	Random Forest	Anomaly classification, split not stated	84.0%	Paper text read
Barraz 2025; Gopalakrishnan 2024, as cited in [4]	Cascading ensemble; NASNet + LSTM	Not stated	88%; 85%	As cited
Millendorf et al. 2020 [1]	Dataset paper	-	No model results	Full paper read
Ours: pretrained cascade	**EfficientNet-B0 at each of 5 stages**	**12 classes, fixed 70/15/15 hold-out**	**83.2%, macro-F1 0.700**	**Our results**
Ours: from-scratch cascade	Five small CNNs	Same hold-out	79.8%, macro-F1 0.646	Our results
Ours: Stage 1 only (pretrained)	Binary triage	Healthy vs defect, same hold-out	95.8%	Our results
Ours: flat CNN / Random Forest	EfficientNet-B0 / raw pixels	12 classes, same hold-out	82.6% / 61.6%	Our results
Ours: Duranay-style baseline	EfficientNet-B0 features + SVM	12 classes, same hold-out	78.7%, macro-F1 0.640	Our results

- **Binary triage.** Our pretrained Stage 1 reaches 95.8% (94.1% from scratch), above the 92.5% of Alves et al. and below the 97.32% of a multi-scale transfer-learning CNN cited by Duranay. It is a single hold-out run, so we claim to be at or slightly above the 92.5% to 93% range and not a new best.
- **12 classes.** Our best, 83.2%, is below the published 12-class figures (84.0%, 85% to 88% for ensembles, 93.93%). We did not reach the 93% ceiling or the 95% target stated in the proposal. A Duranay-style pipeline run on our split scored only 78.7%, below our pretrained cascade, which suggests that the gap to 93.93% is largely due to the evaluation protocol rather than to a weaker model; this is a suggestion and not a proof.
- **Why the results are not like for like.** Duranay uses 10-fold cross-validation, a network trained end to end on the dataset and a supervised feature selection before the cross-validation, so as described, test images can influence these steps. Our test images never do. A raw-pixel Random Forest scores 61.6% on our split against 84.0% in [4], whose protocol is not stated, so part of the gap remains unexplained.

6.7 Data Overlap Analysis

The benchmark has no module or site identifiers, so near-identical crops of the same module can exist. To measure how much this matters, every test image was compared with every training image by the Euclidean distance between their 24 x 40 pixel vectors. Table 6.12 shows the result. A raw-pixel nearest-neighbour classifier with no training already reaches 62.4% accuracy, the same as the raw-pixel Random Forest (61.6%).

Table 6.12: Similarity between test and training images (raw pixels)

Measure	Value
Test images with a training image within distance 0.5	7.4%
Label agreement for those images	95.9%
Test images with a training image within distance 1.0	53.5%
Label agreement for those images	78.1%
Median distance to the nearest training image	0.96
Median distance between random image pairs	4.70

About 7% of the test images have a near-identical training image and about half have a close one, so random fold splits, as used in cross-validated studies, would place such near-duplicates on both sides and raise their scores. This explains part of the difference to published numbers but not all of it, since a Random Forest scores far lower on our split than the 84.0% reported elsewhere.

6.8 Discussion and Limitations

6.8.1 Interpretation of Findings

Three results stand out. Thermal-aware augmentation mattered more than any architectural change. Pretraining gave a large and cheap gain even on 24 x 40 pixel thermal images. And the central hypothesis, that a cascade clearly beats a flat classifier, was only partly supported: the family decision is the weak link, and with equal backbones the two designs tie. The cascade does give something a flat model cannot: it tells where an error came from, and it shows that Stage 1 is a reliable gatekeeper.

6.8.2 Limitations

- **Single runs.** Every configuration was run once with seed 42; differences below about 0.5 points are within noise and no confidence intervals were computed. A 10-fold cross-validation of the pretrained cascade was prepared but did not complete, so it is not reported.

- **Near-duplicate images.** The benchmark has no module or site identifiers, so near-identical crops can fall on both sides of any split.
- **Rare classes.** Soiling, Hot-Spot and Diode-Multi remain weak, which keeps macro-F1 well below accuracy.
- **Efficiency.** Latency was not measured reliably, because timings were taken while the GPU was busy, and the pretrained cascade has five times the parameters of the flat model, so the early exit did not give a demonstrated efficiency advantage.
- **Comparison with published work.** Protocols differ, and some sources could only be checked through abstracts or citing papers.

6.8.3 Threats to Validity

Internal validity is limited by the single seed and by the fact that hyper-parameters were tuned on one validation split. External validity is limited to one dataset of small crops from one provider; the results may not transfer to higher-resolution images or other sites. Construct validity is limited by label noise, which cannot be assessed because no second annotation exists.

6.9 Chapter Summary

The best cascade reaches 83.2% accuracy with ImageNet pretraining and 79.8% from scratch, with Stage 1 at 95.8%. The cascade matches but does not clearly beat a flat classifier, the family stage is its largest error source, and published 12-class results are higher under protocols that are not comparable with a fixed hold-out.

CHAPTER 7

CONCLUSION AND FUTURE SCOPE

7.1 Conclusion

This project implemented Cascaded-SolarNet, a three-stage cascade of five CNNs for 12-class infrared solar module defect classification, with a reproducible split, tests, architecture diagrams and live demonstrations. After removing harmful photometric augmentation, enlarging the input and strengthening the front stages, accuracy rose from 61.7% to 79.8% from scratch and to 83.2% with ImageNet pretraining (macro-F1 0.37 to 0.70). Healthy-versus-defect triage reaches 95.8%, comparable to the 92.5% binary figure published for this dataset.

The claim that a cascade clearly beats a flat classifier was only partly supported. From scratch the cascade is ahead on accuracy and kappa; with the same pretrained backbone it ties a flat CNN while using five times the parameters. Error propagation is the central trade-off: a wrong decision at Stage 1 or Stage 2 cannot be undone, and the family decision remains the largest error source. Published 12-class results (84% to 93.93%) are higher than ours, but they use cross-validation and pretrained features, and the 93% ceiling and 95% target of the proposal were not reached under a stricter fixed hold-out. A practical lesson is that thermal images need thermal-aware augmentation: flips help, while brightness, contrast and noise jitter hurt.

7.2 Future Scope

- **Like-for-like evaluation.** Run several seeds and a leak-free 10-fold cross-validation, and replicate the full NCA pipeline, for a fair comparison with published results.
- **Better family routing.** Train Stage 2 jointly with the specialists or on the cascade's own routing errors, or use soft hand-offs trained end to end.
- **Rare classes.** Apply class-specific augmentation or synthetic oversampling for Soiling, Hot-Spot and Diode-Multi, and deduplicate near-identical crops before splitting.
- **Edge deployment.** Measure real latency, memory and energy on a drone-class device, and test distillation of the pretrained cascade into smaller networks.
- **Additional modalities.** Investigate RGB or contextual information from neighbouring panels, which the crop-level dataset cannot provide.

REFERENCES

[1] M. Millendorf, E. Obropta, and N. Vadhavkar, "Infrared Solar Module Dataset for Anomaly Detection," ICLR 2020 Workshop on AI for Earth Sciences, 2020.

[2] Z. B. Duranay, "Fault Detection in Solar Energy Systems: A Deep Learning Approach," Electronics, vol. 12, no. 21, p. 4397, 2023, doi: 10.3390/electronics12214397.

[3] R. H. F. Alves, G. A. de Deus Junior, E. G. Marra, and R. P. Lemos, "Automatic fault classification in photovoltaic modules using Convolutional Neural Networks," Renewable Energy, vol. 179, pp. 502-516, 2021.

[4] Corley et al., "Aerial Infrared Health Monitoring of Solar Photovoltaic Farms at Scale," CVPR Workshops (PBVS), 2025, arXiv:2503.02128.

[5] S. P. Pathak and S. A. Patil, "Enhanced Fault Identification in Solar Panels through Binary Cascaded Convolutional Classifiers with Thermal-Visual Augmentation," Int. J. Image, Graphics and Signal Processing, vol. 17, no. 4, 2025.

[6] "PV-faultNet: Optimized Convolutional Neural Network to Detect Defects in PV Production," arXiv:2411.02997, 2024.

[7] "Defect Detection of Photovoltaic Modules Based on Improved VarifocalNet," Scientific Reports, vol. 14, 2024, doi: 10.1038/s41598-024-66234-3.

[8] U. Otamendi, I. Martinez, M. Quartulli, et al., "Segmentation of Cell-Level Anomalies in Electroluminescence Images of Photovoltaic Modules," Solar Energy, vol. 220, pp. 914-926, 2021.

[9] H. Tella, A. Hussein, S. Rehman, et al., "PV Panel Defect Classification Using Deep Learning Ensemble Methods," Case Studies in Thermal Engineering, vol. 66, 2025.

[10] Z. Cai and N. Vasconcelos, "Cascade R-CNN: Delving into High Quality Object Detection," in Proc. CVPR, 2018, pp. 6154-6162.

[11] M. Tan and Q. V. Le, "EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks," in Proc. ICML, 2019, arXiv:1905.11946.

[12] K. He, X. Zhang, S. Ren, and J. Sun, "Deep Residual Learning for Image Recognition," in Proc. CVPR, 2016, pp. 770-778.

[13] A. Paszke et al., "PyTorch: An Imperative Style, High-Performance Deep Learning

Library," in Advances in Neural Information Processing Systems 32, 2019.

[14] J. Deng et al., "ImageNet: A Large-Scale Hierarchical Image Database," in Proc. CVPR, 2009.

[15] Z. Barraz, I. Sebari, N. Lamrini, K. Ait El Kadi, and I. Ait Abdelmoula, "A cascading decision system for enhanced anomaly classification of large-scale photovoltaic systems using drone's thermal data with class-imbalance problem," Results in Engineering, vol. 25, p. 103876, 2025.

[16] S. Gopalakrishnan, N. I. Abdul Wahab, V. Veerasamy, H. Hizam, and R. A. Farade, "NASNet-LSTM based deep learning classifier for anomaly detection in solar photovoltaic modules," Journal of Physics: Conference Series, p. 012006, IOP Publishing, 2024.

APPENDIX A

CONFIGURATION FILES

This appendix reproduces the two configuration files that define the best from-scratch experiment and the pretrained experiment. Every reported result can be regenerated from the configuration, the saved split file and the seed.

Listing A.1: configs/improved3.yaml (from-scratch cascade)

```
# improved3 = improved with larger triage (CNN-1) and family (CNN-2) networks.
# Fine CNNs are copied from runs_improved; only triage and family are retrained.
# Improved: 80x48 upsampling, two convs per block, no photometric jitter, longer
schedule.
data:
  root: data/InfraredSolarModules     # folder containing module_metadata.json
  splits_file: data/splits.json       # written by scripts/make_splits.py
  seed: 42
  val_frac: 0.15
  test_frac: 0.15
  upsample: 2                         # 1 = native 40x24; 2 = 80x48 (optional, see
report 8.5)

augment:
  hflip: 0.5
  vflip: 0.5
  brightness: 0.0
  contrast: 0.0
  noise_std: 0.0

train:
  batch_size: 64
  epochs: 150
  lr: 5.0e-4
  weight_decay: 1.0e-4
  patience: 25                        # early stopping on val macro-F1
  weighted_loss: true                 # w_r = N / (K * n_r)
  num_workers: 2
  out_dir: runs_improved3

# Per-stage CNN widths (poster, "Class Hierarchy" panel).
# Each entry is a Conv3x3-BN-ReLU-MaxPool block; head is GlobalAvgPool + FC.
models:
  triage:     {channels: [64, 128, 128],  dropout: 0.2, convs_per_block: 2}   # CNN-1:
No-Anomaly vs Defect
  family:     {channels: [128, 256, 128], dropout: 0.4, convs_per_block: 2}   # CNN-2:
Electrical / Environmental / Structural
  fine:       {channels: [128, 256, 128], dropout: 0.4, convs_per_block: 2}  # CNN-
3a/b/c: one per family
  flat:       {channels: [128, 256, 128], dropout: 0.4, convs_per_block: 2}  # single-
stage 12-class baseline
```

Listing A.2: configs/pretrained.yaml (pretrained cascade)

```
# Pretrained cascade: an ImageNet-pretrained EfficientNet-B0 under every stage, fine-
tuned.
# Inputs are resized inside the model (to `size`), so data.upsample stays 1.
data:
  root: data/InfraredSolarModules     # folder containing module_metadata.json
  splits_file: data/splits.json       # written by scripts/make_splits.py
  seed: 42
```

```
  val_frac: 0.15
  test_frac: 0.15
  upsample: 1

augment:
  hflip: 0.5
  vflip: 0.5
  brightness: 0.0
  contrast: 0.0
  noise_std: 0.0

train:
  batch_size: 64
  epochs: 25
  lr: 3.0e-4
  weight_decay: 1.0e-4
  patience: 6                          # early stopping on val macro-F1
  weighted_loss: true                  # w_r = N / (K * n_r)
  num_workers: 2
  amp: true                            # mixed precision, ~2x faster on GPU
  out_dir: runs_pretrained

models:
  triage:     {pretrained: true, size: [160, 96], dropout: 0.2}
  family:     {pretrained: true, size: [160, 96], dropout: 0.3}
  fine:       {pretrained: true, size: [160, 96], dropout: 0.3}
  flat:       {pretrained: true, size: [160, 96], dropout: 0.3}
```

Listing A.3: configs/taxonomy.yaml (semantic class hierarchy)

```
# Class hierarchy for Cascaded-SolarNet.
#
# Stage 1 separates `healthy` from everything else (early exit).
# Stage 2 routes a defect to one of the families below.
# Stage 3 runs that family's fine-grained CNN.
#
# Label strings must match `anomaly_class` in InfraredSolarModules'
# module_metadata.json exactly. Regroup classes here; no code changes needed.

healthy: No-Anomaly

families:
  Electrical:
    - Cell
    - Cell-Multi
    - Diode
    - Diode-Multi
    - Hot-Spot
    - Hot-Spot-Multi
  Environmental:
    - Soiling
    - Shadowing
    - Vegetation
  Structural:
    - Cracking
    - Offline-Module
```

APPENDIX B

KEY SOURCE CODE

The listings show the three pieces of code that define the behaviour of the system: the stage network and the cascade routing, the augmentation, and the attribution of errors to stages. The complete code is available in the project repository.

Listing B.1: Stage network and cascade routing (src/solarnet/models.py)

```
class StageCNN(nn.Module):
    """[Conv3x3-BN-ReLU-MaxPool2] x len(channels) -> GlobalAvgPool -> Dropout -> FC.

    With the native 40x24 input, three blocks leave a 5x3 map before pooling.
    """

    def __init__(self, channels: list[int], num_classes: int, dropout: float = 0.0,
in_ch: int = 1,
                 convs_per_block: int = 1):
        super().__init__()
        layers, prev = [], in_ch
        for ch in channels:
            for _ in range(convs_per_block):
                layers += [
                    nn.Conv2d(prev, ch, 3, padding=1, bias=False),
                    nn.BatchNorm2d(ch),
                    nn.ReLU(inplace=True),
                ]
                prev = ch
            layers.append(nn.MaxPool2d(2))
        self.features = nn.Sequential(*layers)
        self.head = nn.Sequential(nn.Flatten(), nn.Dropout(dropout), nn.Linear(prev,
num_classes))
        for m in self.modules():
            if isinstance(m, nn.Conv2d):
                nn.init.kaiming_normal_(m.weight, mode="fan_out", nonlinearity="relu")

    def forward(self, x: torch.Tensor) -> torch.Tensor:
        x = self.features(x)
        return self.head(F.adaptive_avg_pool2d(x, 1))
```

Listing B.2: Cascade inference with early exit (src/solarnet/models.py)

```
    def predict(self, x: torch.Tensor, threshold: float = 0.5, tta: bool = False) ->
dict[str, torch.Tensor]:
        """Route a batch through the cascade (hard routing, with early exit).

        `threshold` is the triage P(defect) above which an image continues to stage 2.
        Returns global class predictions plus the per-stage decisions, so
        evaluation can attribute errors to the stage that made them.
        """
        n = x.shape[0]
        pred = torch.full((n,), self.healthy_index, dtype=torch.long, device=x.device)
        fam_pred = torch.full((n,), -1, dtype=torch.long, device=x.device)

        is_defect = self._probs(self.triage, x, tta)[:, 1] > threshold
        idx = is_defect.nonzero(as_tuple=True)[0]
        if idx.numel():  # early exit: healthy panels never reach stage 2
            fam_pred[idx] = self._probs(self.family, x[idx], tta).argmax(1)
            for f_i, fam in enumerate(self.family_names):
```

```
                sub = idx[fam_pred[idx] == f_i]
                if sub.numel():
                    pred[sub] = self.class_offsets[fam] + self._probs(self.fine[fam],
x[sub], tta).argmax(1)
        return {"pred": pred, "is_defect": is_defect, "family": fam_pred}

    @torch.no_grad()
```

Listing B.3: Training augmentation (src/solarnet/common.py)

```
def augment(x: torch.Tensor, aug: dict) -> torch.Tensor:
    """Stochastic train-time augmentation on a [B,1,H,W] batch in [0,1].

    90-degree rotations from the data report are left out: the images are
    40x24, so a quarter turn would change the shape. Two flips together
    already give the 180-degree rotation.
    """
    b = x.shape[0]
    flip_h = torch.rand(b, device=x.device) < aug["hflip"]
    x = torch.where(flip_h[:, None, None, None], x.flip(-1), x)
    flip_v = torch.rand(b, device=x.device) < aug["vflip"]
    x = torch.where(flip_v[:, None, None, None], x.flip(-2), x)

    def jitter(scale: float) -> torch.Tensor:
        return 1 + (torch.rand(b, 1, 1, 1, device=x.device) * 2 - 1) * scale

    x = x * jitter(aug["brightness"])
    mean = x.mean(dim=(1, 2, 3), keepdim=True)
    x = (x - mean) * jitter(aug["contrast"]) + mean
    x = x + torch.randn_like(x) * aug["noise_std"]
    return x.clamp(0, 1)
```

Listing B.4: Attribution of errors to stages (src/solarnet/evaluate.py)

```
def stage_errors(y: np.ndarray, out: dict, tax: Taxonomy) -> dict:
    """Attribute every cascade mistake to the first stage that went wrong."""
    healthy = tax.index(tax.healthy)
    true_defect = y != healthy
    true_fam = np.array([tax.family_names.index(tax.family_of(tax.classes[c])) if c !=
healthy else -1 for c in y])
    is_defect = out["is_defect"].numpy()
    fam = out["family"].numpy()
    wrong = out["pred"].numpy() != y

    triage_miss = true_defect & ~is_defect           # defect exited early as healthy
    triage_false_alarm = ~true_defect & is_defect    # healthy sent down the cascade
    family_wrong = true_defect & is_defect & (fam != true_fam)
    fine_wrong = wrong & ~triage_miss & ~triage_false_alarm & ~family_wrong
    return {
        "stage1_missed_defects": int(triage_miss.sum()),
        "stage1_false_alarms": int(triage_false_alarm.sum()),
        "stage2_wrong_family": int(family_wrong.sum()),
        "stage3_wrong_subtype": int(fine_wrong.sum()),
        "total_errors": int(wrong.sum()),
        "early_exit_rate": float((~is_defect).mean()),
    }
```

Listing B.5: Class weights for the imbalanced loss (src/solarnet/data.py)

```
def class_weights(targets: np.ndarray, num_classes: int) -> np.ndarray:
    """w_r = N / (K * n_r); classes absent from `targets` get weight 0."""
    counts = np.bincount(targets, minlength=num_classes).astype(np.float64)
```

```
    with np.errstate(divide="ignore"):
        w = np.where(counts > 0, len(targets) / (num_classes * counts), 0.0)
    return w.astype(np.float32)
```

APPENDIX C

RUNNING THE DEMONSTRATIONS

Two demonstrations are provided. Both need the trained checkpoints and the dataset, which is downloaded by a script from the public repository of the dataset.

- **Command-line demonstration.** python scripts/predict.py --sample 5 draws, for random test images, the decision and confidence of each stage and saves the figure. The option --class Hot-Spot restricts the sample to one class and --image path loads any image.
- **Browser demonstration.** python scripts/make_demo_html.py creates a single HTML page with three test images per class. Opening it in a browser shows the selected image, and pressing Run cascade reveals Stage 1, Stage 2 and Stage 3 one after another, with the final prediction in green if it is correct and red if not.
- **Reproducing the results.** python scripts/make_splits.py creates the split, python -m solarnet.train --node all --config configs/pretrained.yaml trains the five networks, and python -m solarnet.evaluate --flat --config configs/pretrained.yaml evaluates the cascade and the flat baseline.

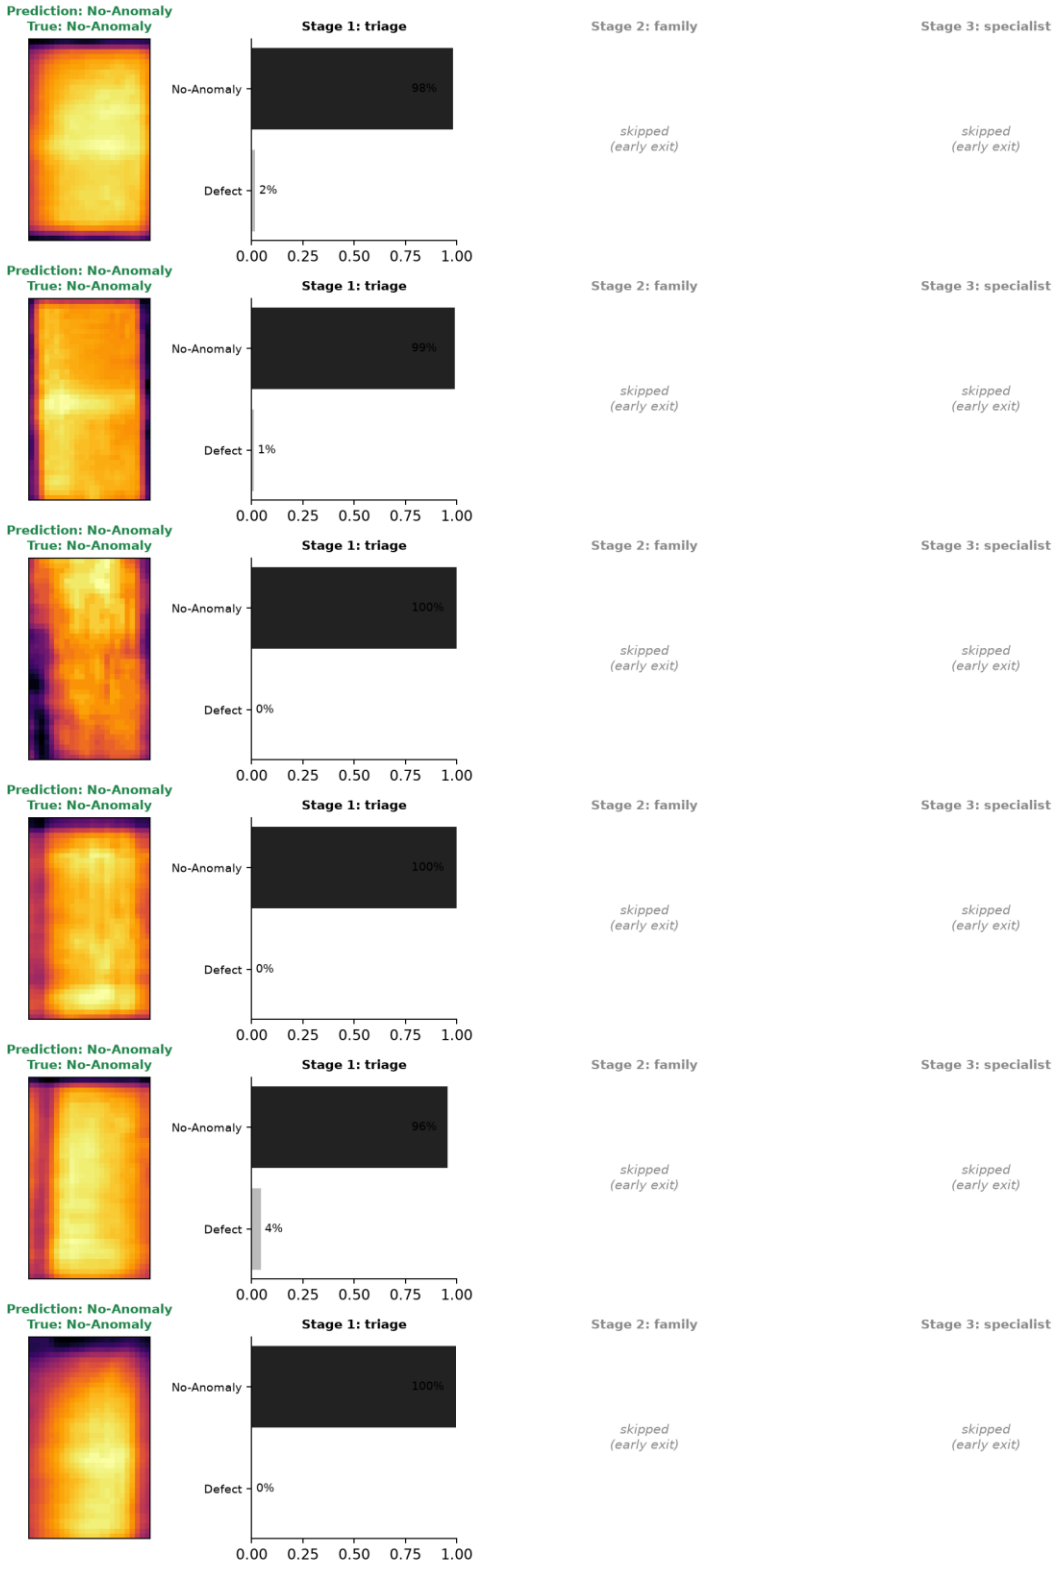


Fig. 15. Example output of the command-line demonstration for random test images (green: correct, red: wrong).

APPENDIX D

REPOSITORY STRUCTURE

Table D.1: Layout of the project repository

Path	Content
src/solarnet/	Package: taxonomy, data, models, common helpers, training, evaluation
configs/	Experiment configurations and the class hierarchies (semantic and data-driven)
scripts/	Download and split scripts, baselines, cross-validation, figure and report generators, demonstrations
tests/	Automated tests that run on synthetic data
results/	Raw metrics of every reported experiment (JSON)
checkpoints/	Trained Stage 3 specialists shared with the team
docs/	Architecture diagrams and the reports
notebooks/	Notebook that runs the whole pipeline on Colab